

Star-Wars® Jedi Workshop

Version 1.5

Mark Quinsland
Senior Field Engineer, Neo4j
March 13, 2023



Neo4j Star-Wars® Jedi Workshop Agenda

- Quick review of Neo4j
- Creating a Jedi data model using Arrows
- Using AuraDB - Neo4j's Cloud DB
- Creating a Star-Wars® graph using Importer
- Exploring the graph visually using Bloom
- Cypher Queries to read, write, update and delete
- Loading Jedi data into Neo4j using Python
- Creating Jedi DataFrames using Jupyter notebooks
- Analyzing characters using Graph Data Science algorithms

The World's Most Popular Property Graph DB

100m+ Neo4j Downloads
250k+ Community Members

7 Retailers
of the Top 10



7 Telcos
of the Top 10



8 Automakers
of the Top 10



8 Insurance
of the Top 10



5 Pharmaceuticals
of the Top 5



20 Banks
of the Top 20
North American



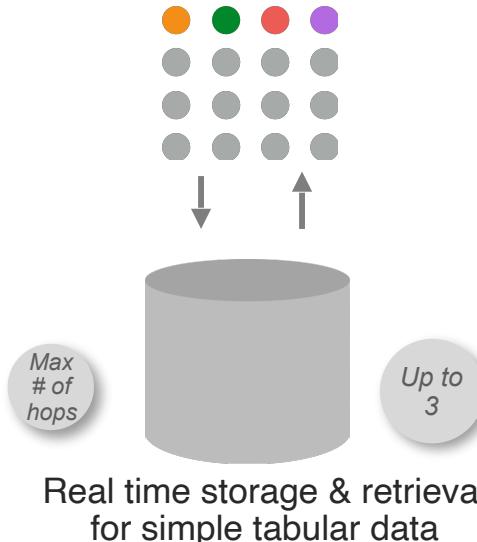
3 Aircraft
Manufacturers
of the Top 5



Data Platform Evolution

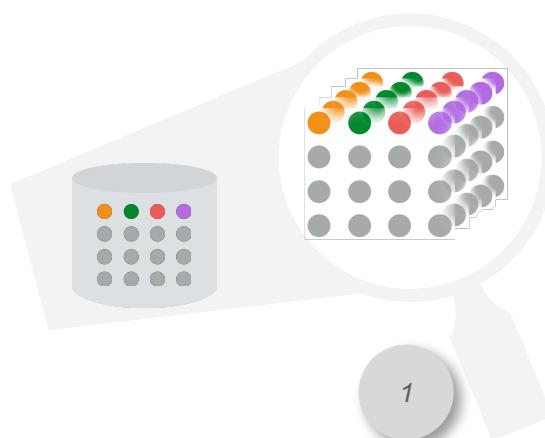
TRADITIONAL DATABASES

Store and retrieve data

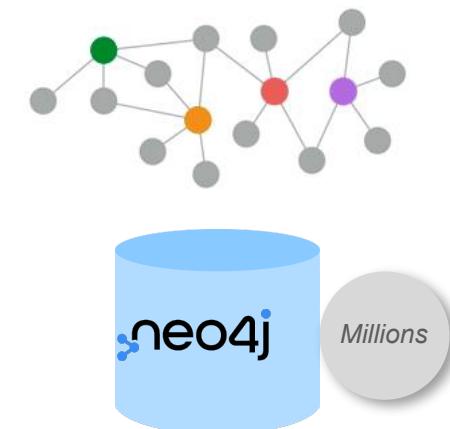


BIG DATA TECHNOLOGY

Aggregate and filter data



Exploit connections in data



Real-time graph queries & algorithms for contextual insights

"Our **Neo4j** solution is literally **thousands of times faster** than the prior MySQL solution, with queries that require **10-100 times less code**"
Volker Pacher, Senior Developer



the Labeled Property Graph Model

Nodes

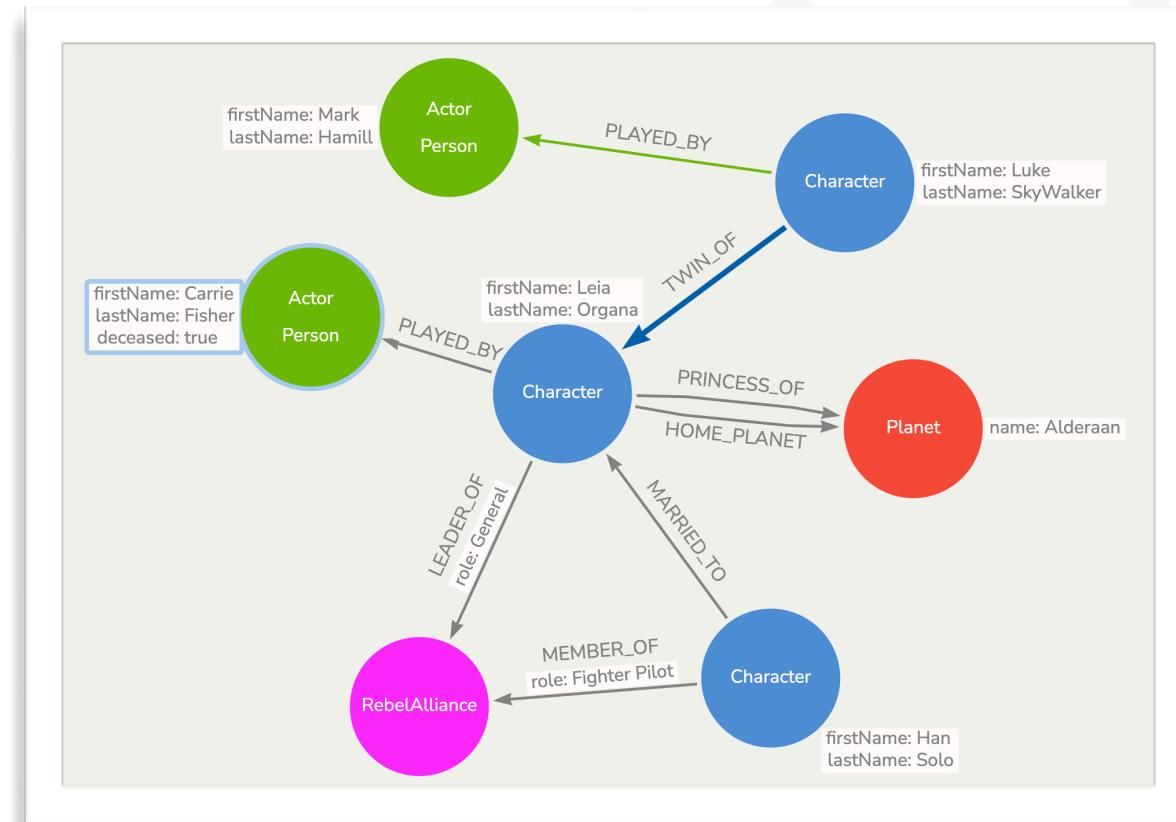
- Define the nouns/entities
- Can have multiple labels

Relationships

- Verbs that relate nodes by type and direction
- Nodes can have multiple relationships

Properties

- Attributes of Nodes and Relationships
- Stored as name/value pairs



- New labels, relationship types and properties can be added w/o schema change



Complex SQL Queries vs Cypher

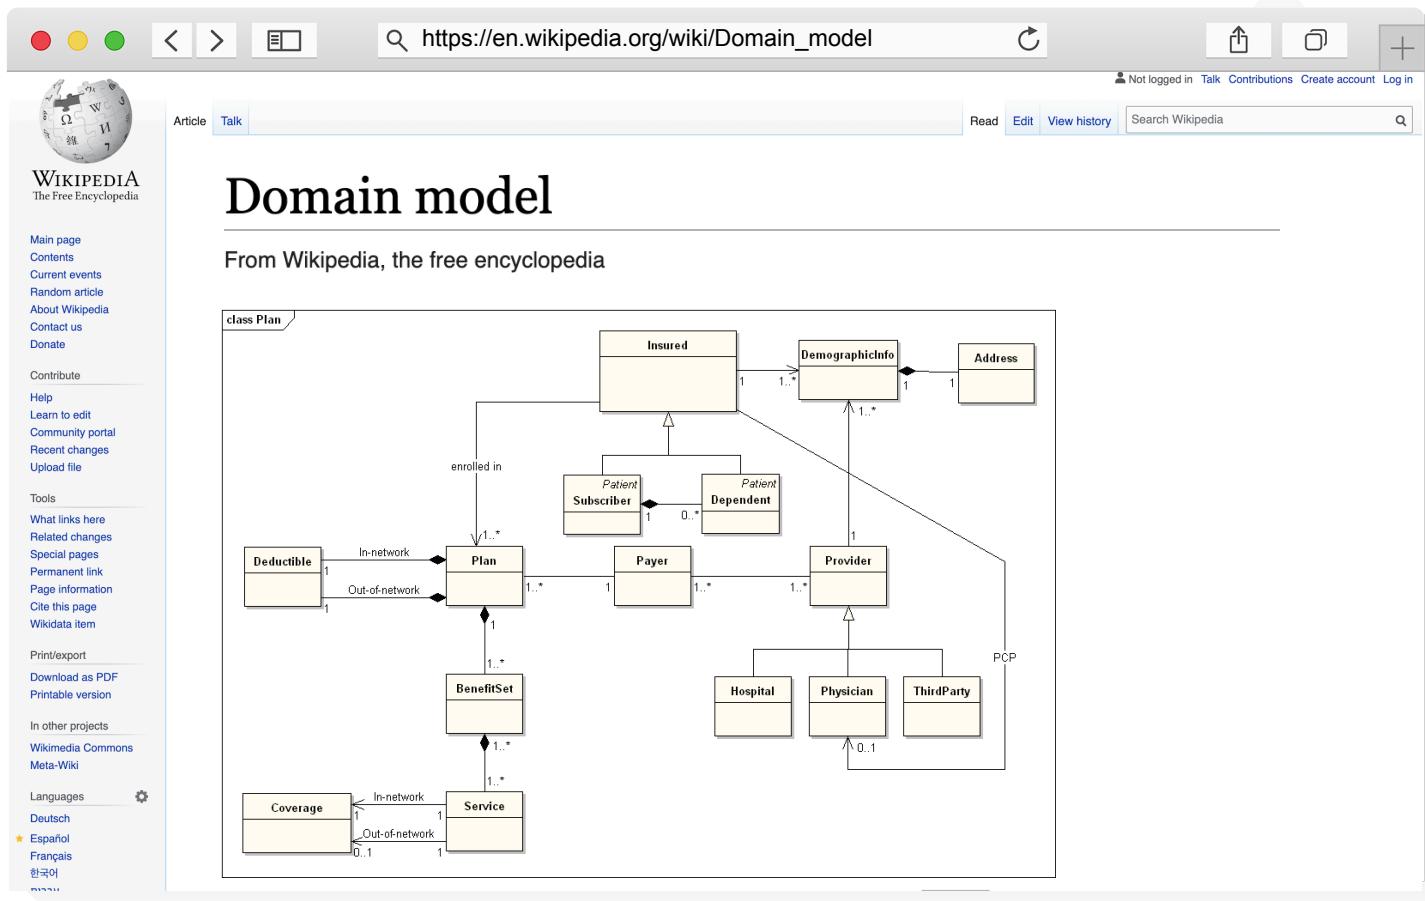
Find all direct reports and how many people they manage, up to three levels down

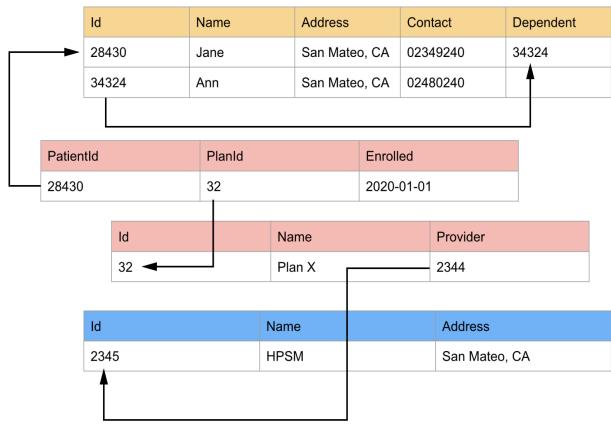
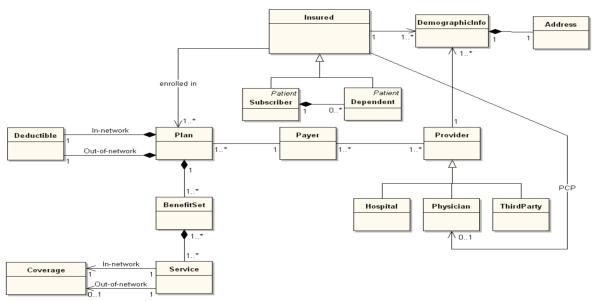
```
(SELECT T.directReportees AS directReportees, sum(T.count) AS count
FROM (
  SELECT manager.pid AS directReportees, 0 AS count
  FROM person_reportee_manager
  WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
UNION
  SELECT manager.pid AS directReportees, count(manager.directly_manages) AS count
  FROM person_reportee_manager
  WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
  GROUP BY directReportees
UNION
  SELECT manager.pid AS directReportees, count(reportee.directly_manages) AS count
  FROM person_reportee_manager
  JOIN person_reportee_reportee
  ON manager.directly_manages = reportee.pid
  WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
  GROUP BY directReportees
UNION
  SELECT manager.pid AS directReportees, count(L2Reportees.directly_manages) AS count
  FROM person_reportee_manager
  JOIN person_reportee L1Reportees
  ON manager.directly_manages = L1Reportees.pid
  JOIN person_reportee L2Reportees
  ON L1Reportees.directly_manages = L2Reportees.pid
  WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
  GROUP BY directReportees
) AS T
GROUP BY directReportees)
UNION
(SELECT T.directReportees AS directReportees, sum(T.count) AS count
FROM (
  SELECT manager.pid AS directReportees, 0 AS count
  FROM person_reportee_manager
  WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
UNION
  SELECT reportee.pid AS directReportees, count(reportee.directly_manages) AS count
  FROM person_reportee_manager
  JOIN person_reportee_reportee
  ON manager.directly_manages = reportee.pid
  WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
  GROUP BY directReportees
UNION
  SELECT L2Reportees.pid AS directReportees, count(L2Reportees.directly_manages) AS count
  FROM person_reportee_manager
  JOIN person_reportee L1Reportees
  ON manager.directly_manages = L1Reportees.pid
  JOIN person_reportee L2Reportees
  ON L1Reportees.directly_manages = L2Reportees.pid
  WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
  GROUP BY directReportees
) AS T
GROUP BY directReportees)
UNION
(SELECT T.directReportees AS directReportees, sum(T.count) AS count
FROM (
  SELECT manager.pid AS directReportees, 0 AS count
  FROM person_reportee_manager
  WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
UNION
  SELECT L2Reportees.pid AS directReportees, count(L2Reportees.directly_manages) AS count
  FROM person_reportee_manager
  JOIN person_reportee L1Reportees
  ON manager.directly_manages = L1Reportees.pid
  JOIN person_reportee L2Reportees
  ON L1Reportees.directly_manages = L2Reportees.pid
  WHERE manager.pid = (SELECT id FROM person WHERE name = "fName lName")
  GROUP BY directReportees
) AS T
GROUP BY directReportees)
```

Cypher Query

```
MATCH (boss)-[:MANAGES*0..6]->(sub),
      (sub)-[:MANAGES*1..3]->(report)
WHERE boss.name = "John Doe"
RETURN sub.name AS Subordinate,
       count(report) AS Total
```

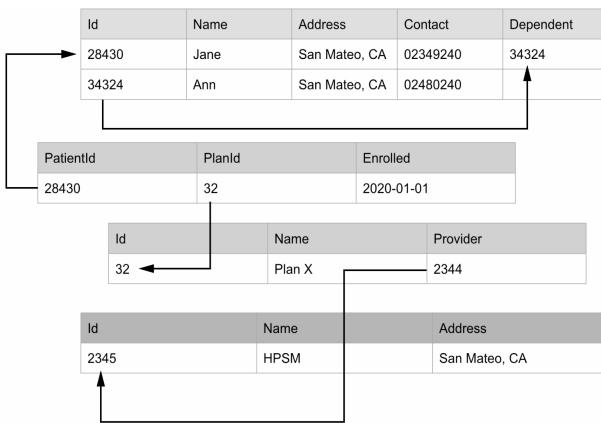
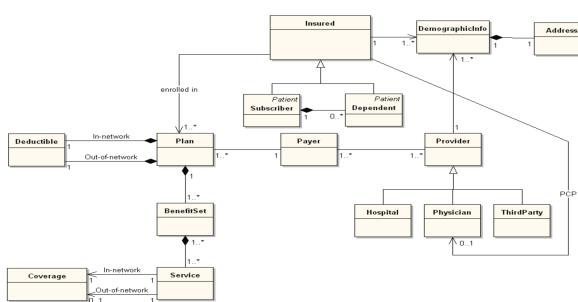
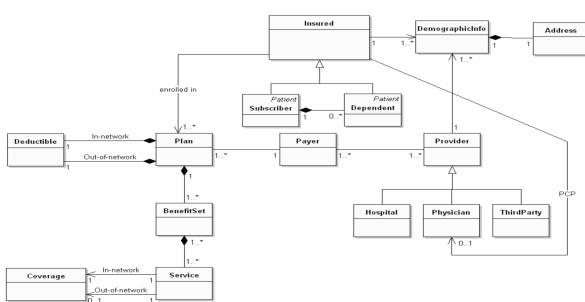
- Less time writing queries
- Less time debugging queries
- Code that's easier to read



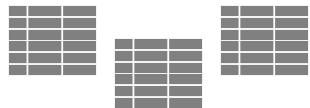


Relational





Relational

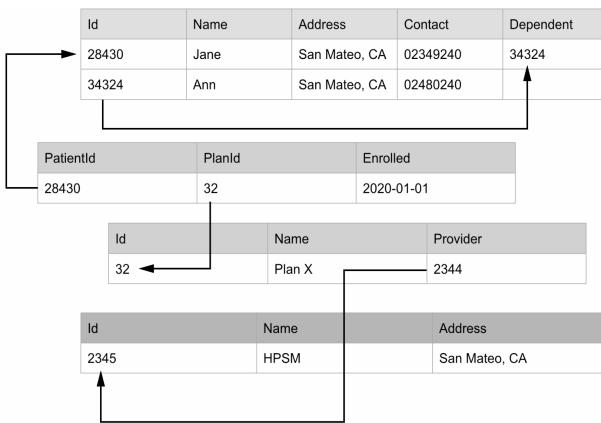
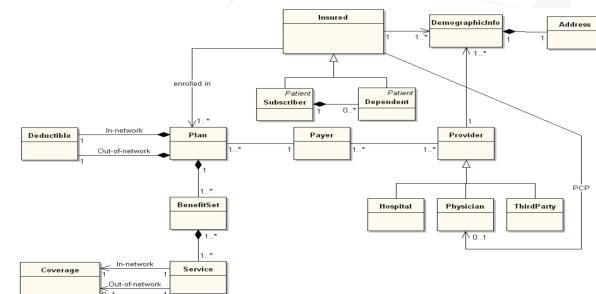
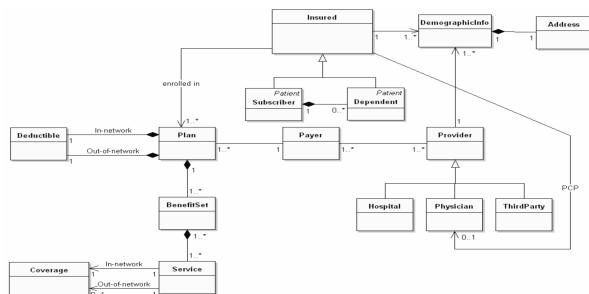
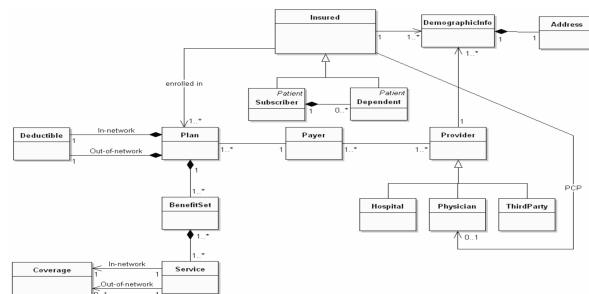


```
// Patient
{
  "_id": 593340651,
  "first": "Gregorio",
  "last": "Lang",
  "addr": {
    "street": "623 Flowers Rd",
    "city": "Groton",
    "state": "NH",
    "zip": 3266
  },
  "physicians": [
    10387 33456
  ],
  "procedures": [
    {"id": "551ac", "type": "Chest X-ray"},
    {"id": "343fs", "type": "Blood Test"}
  ]
}

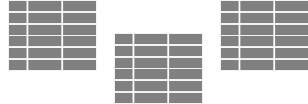
// Procedure
{
  "_id": "551ac",
  "date": "2000-04-26",
  "hospital": 161,
  "patient": 593340651,
  "physician": 10387,
  "type": "Chest X-ray",
  "records": [ "67bc6" ]
}
```

Document





Relational

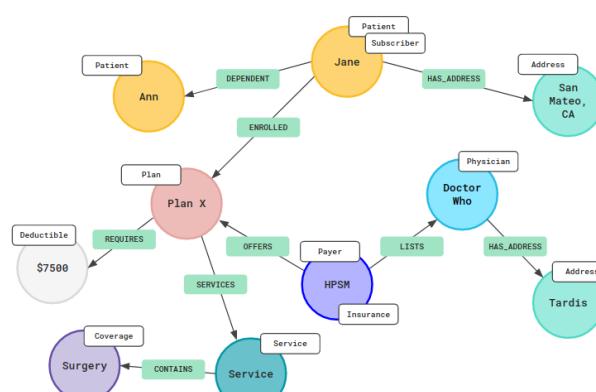


```

// Patient
{
  "_id": "593340651",
  "first": "Gregorio",
  "last": "Lang",
  "addr": {
    "street": "623 Flowers Rd",
    "city": "Groton",
    "state": "NH",
    "zip": 3266
  },
  "physicians": [
    10387 33456
  ],
  "procedures": [
    {"id": "551ac", "type": "Chest X-ray"}, {"id": "343fs", "type": "Blood Test"}
  ]
}

// Procedure
{
  "_id": "551ac",
  "date": "2000-04-26",
  "hospital": 161,
  "patient": 593340651,
  "physician": 10387,
  "type": "Chest X-ray",
  "records": [ "67bc6" ]
}
  
```

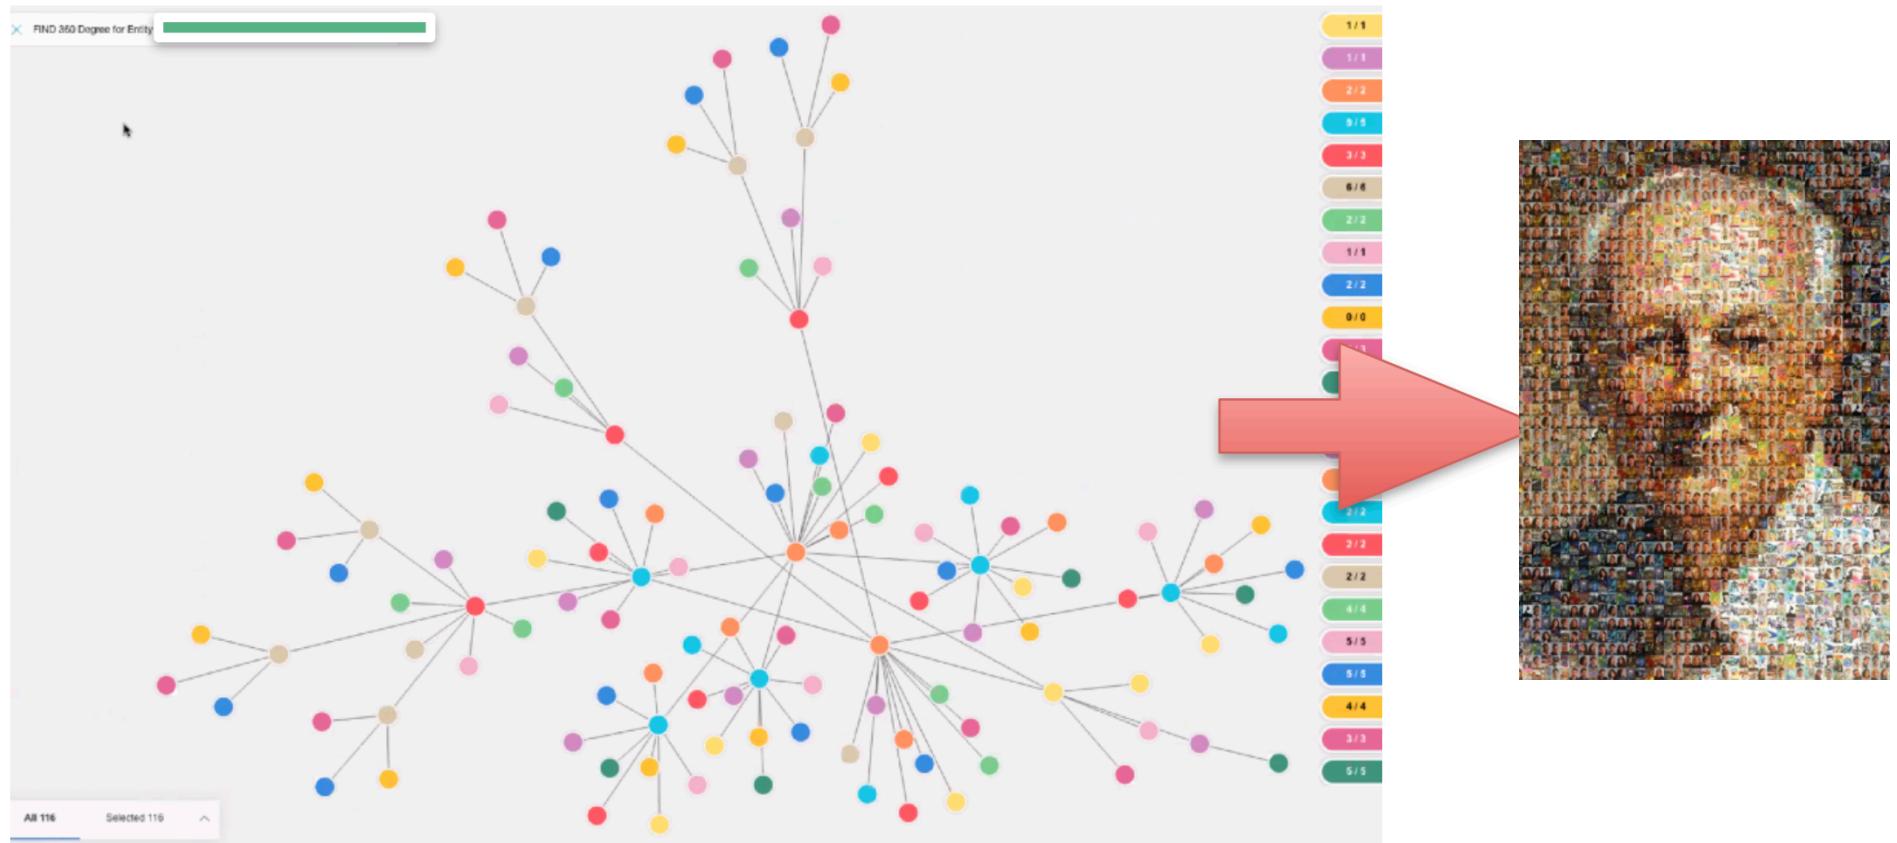
Document



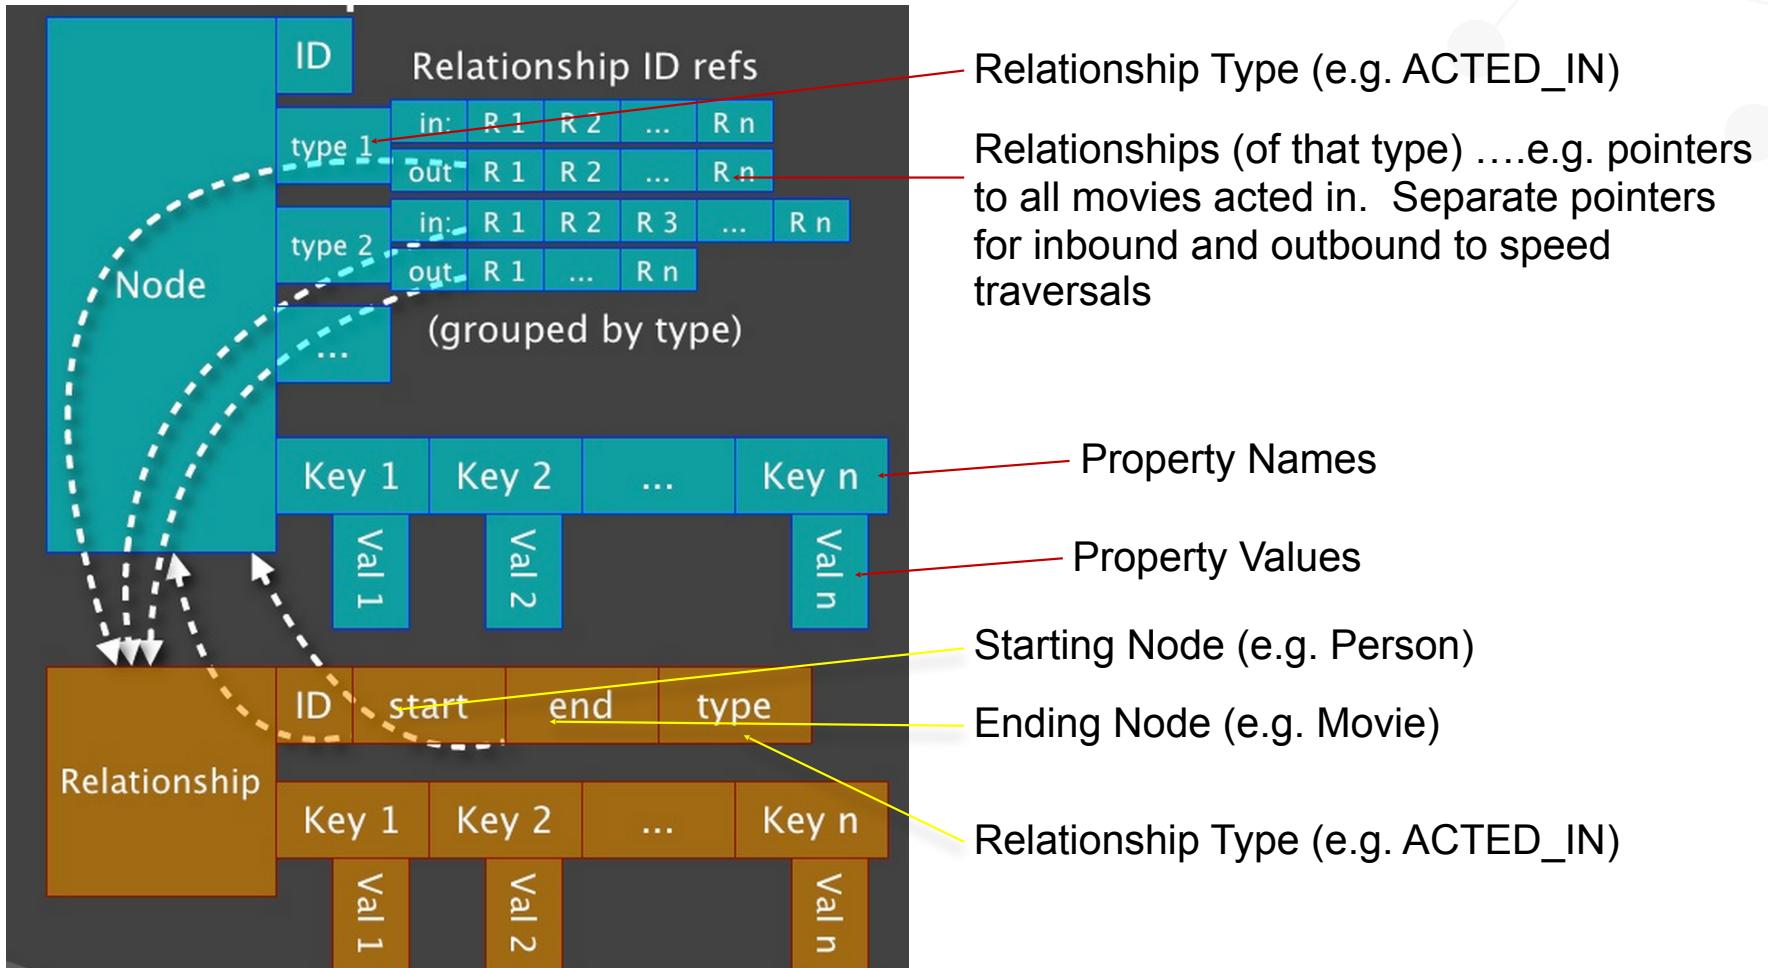
Graph



Real-time Joins of Data From Myriad Sources to Provide 360° View



Neo4j Native Graph Data Storage - Why Neo4j is so Fast



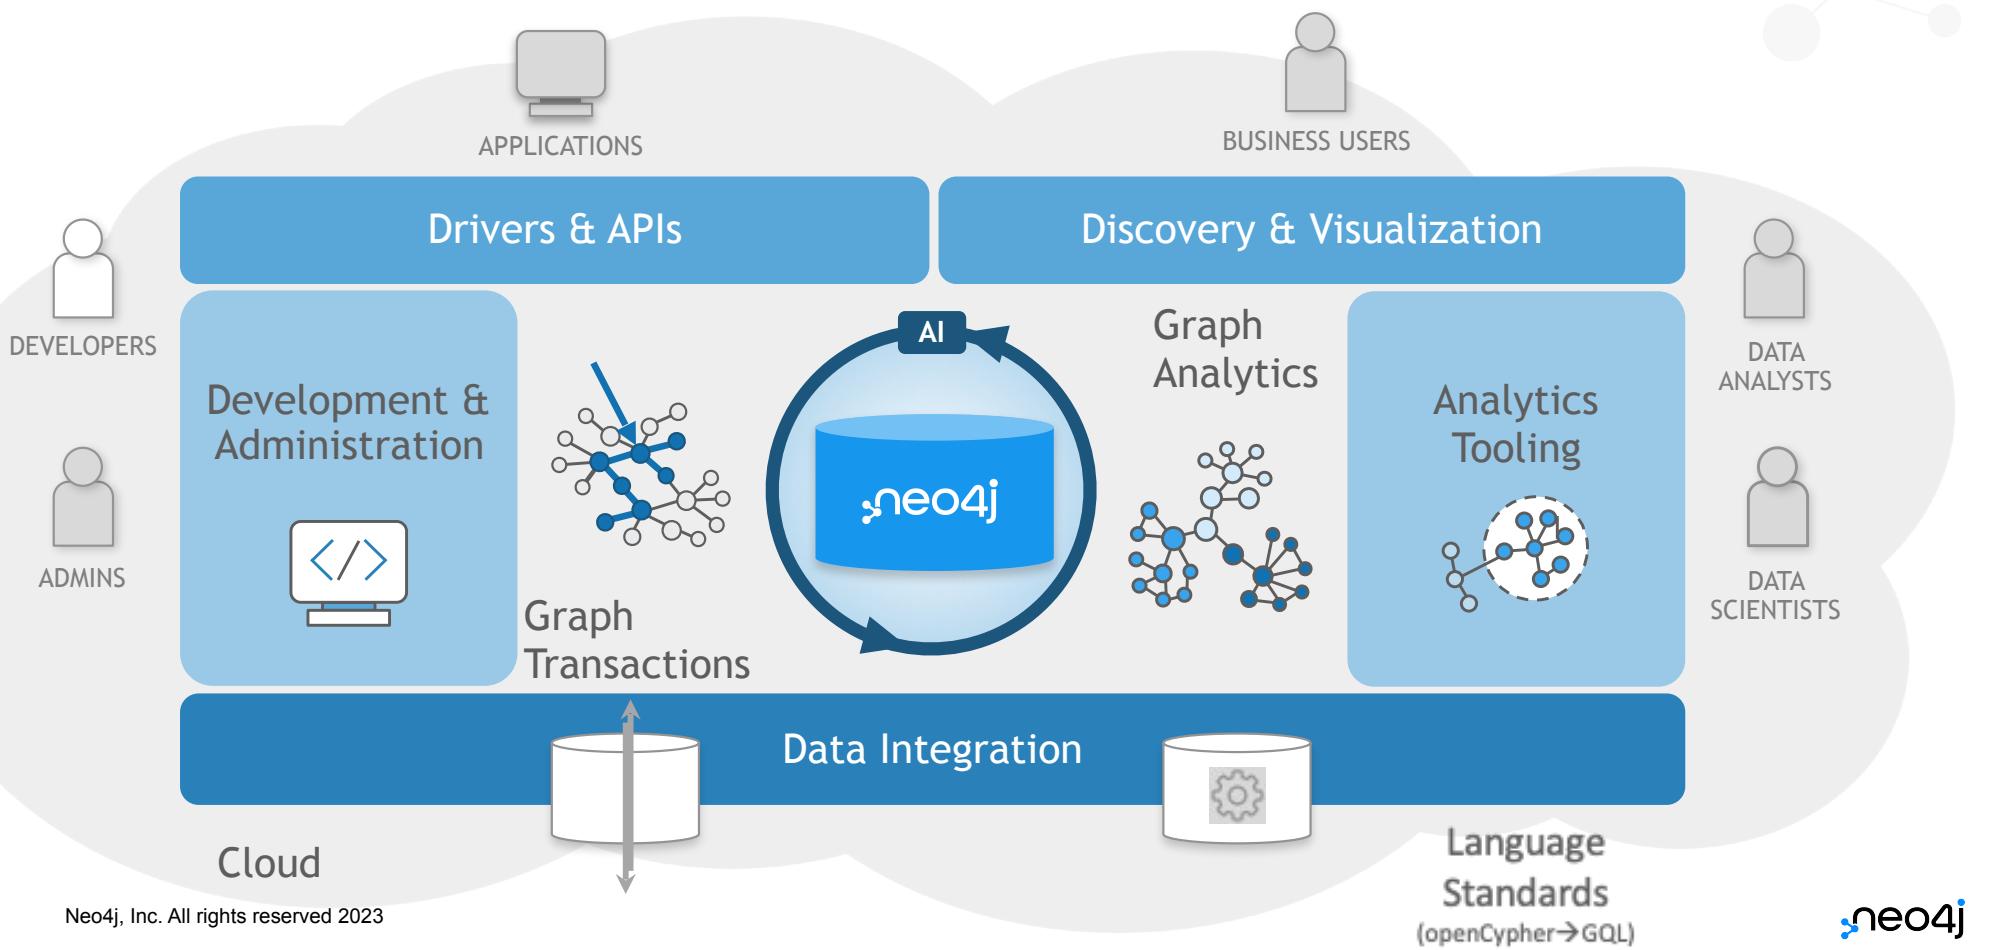
Understanding Latency

System Event	Actual Latency	Relative Latency
One CPU Cycle	0.4 ns	1 second
RAM Access (DDR-DIMM)	~100 ns	4 minutes
SSD I/O	50-150 µs	1.5 - 4 days
Rotational disk I/O	1-10 ms	1 - 10 months
Intra-AZ latency	0.5 ms - 1.0 ms	2 - 4 weeks
Inter-AZ latency	1-10 ms	1 - 10 months
SF to NYC Internet	65 ms	5 years

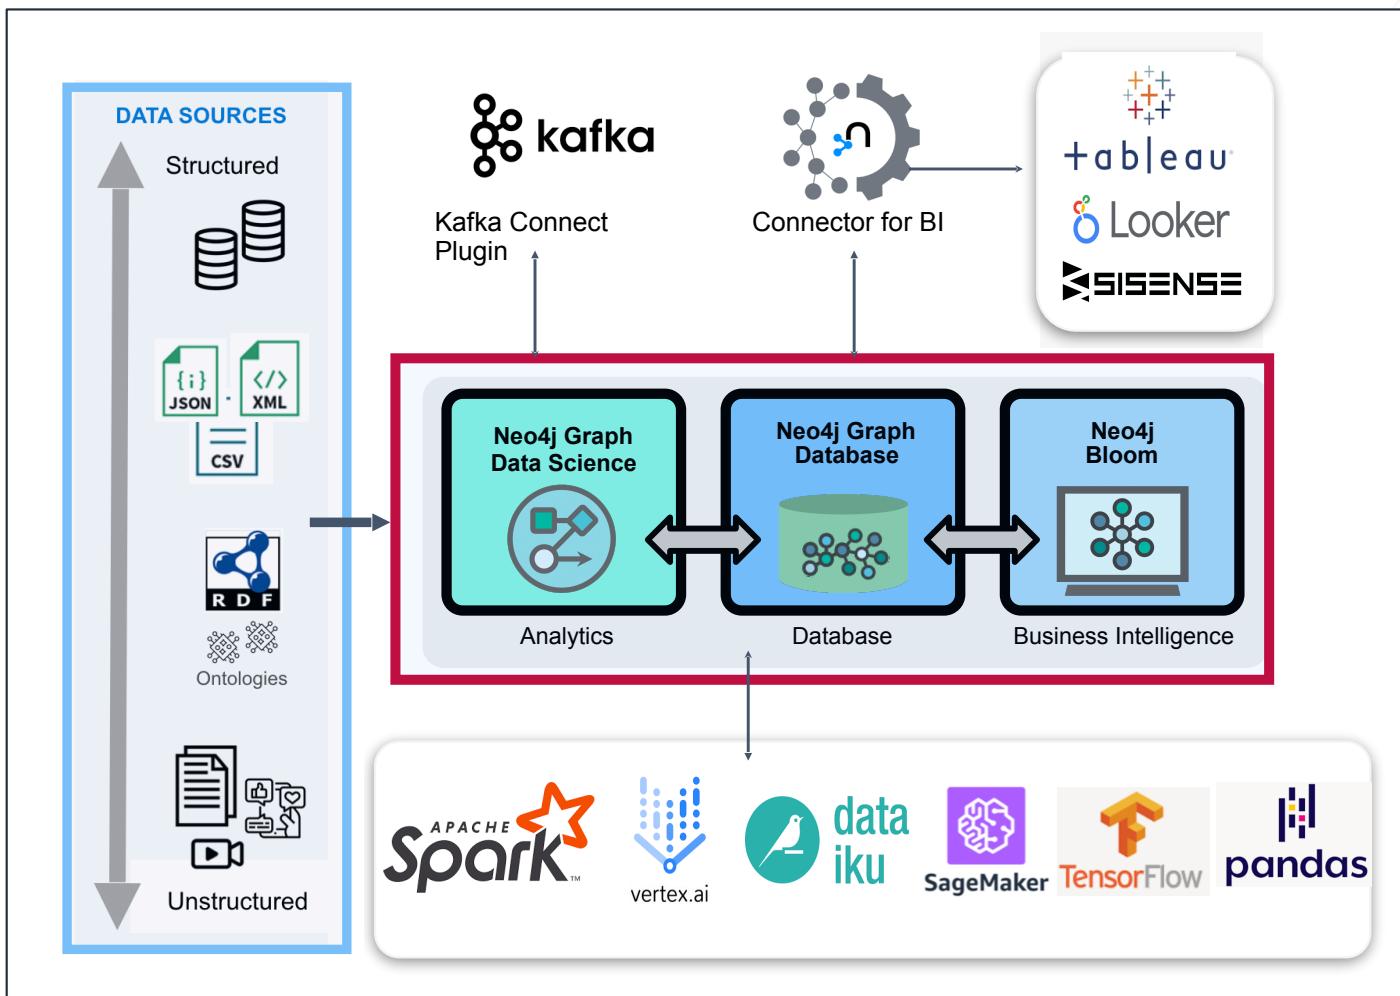
<https://www.prowesscorp.com/computer-latency-at-a-human-scale/>

¹ Relative Latency Using One CPU Cycle = One Second

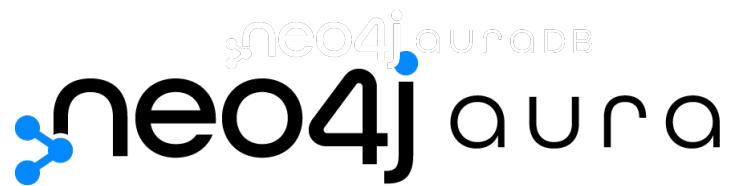
Neo4j Graph Platform



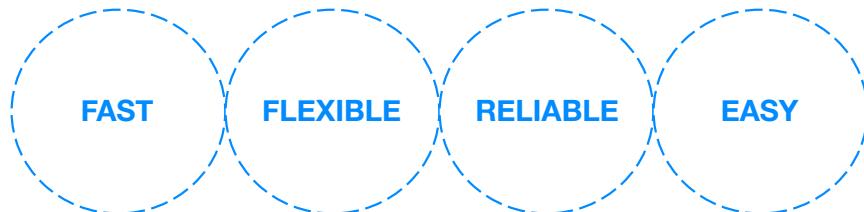
Neo4j Data Science Ecosystem



Neo4j in the Cloud



Fully-managed Cloud Service



Automated Upgrades, Maintenance

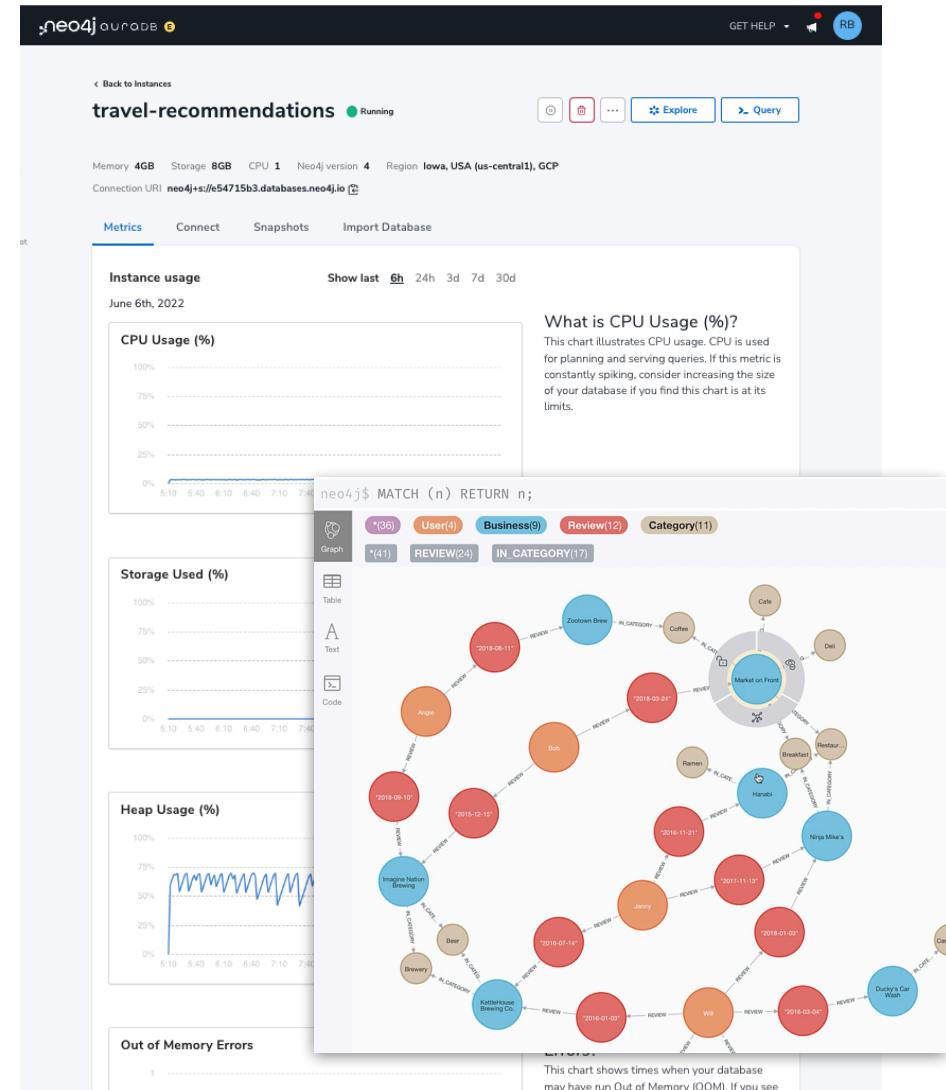
Scalable and Elastic, On-Demand

Enterprise-Grade Security

High Availability

Simple Pricing, Consumption-Based

Procure Direct or via Cloud Marketplace



AuraDB Free

For small development projects, learning, experimentation and prototyping.

- AuraDB is Neo4j's fully managed cloud service
 - Zero-administration
 - Always-on graph database
- It's available today for free, forever
 - limited to 200k nodes and 400k relationships
 - Includes Bloom and APOC
- No credit card is required to register

neo4j.com/cloud/aura/



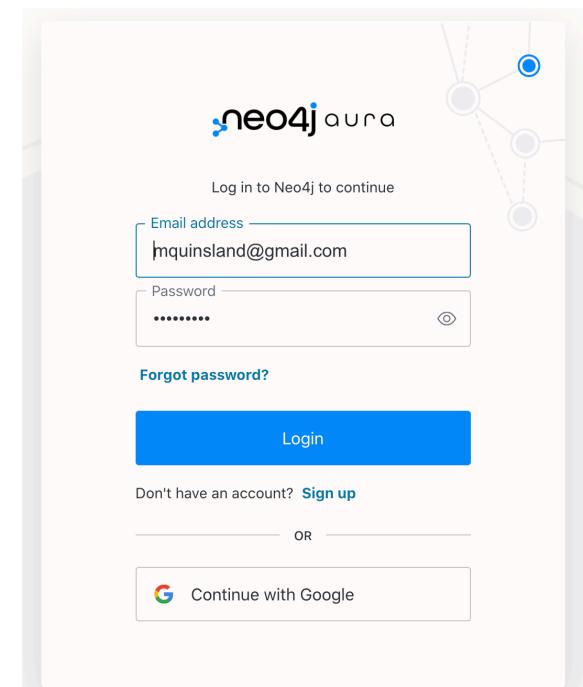
Lab 0 - Connecting to Aura



Lab 0.0 - Creating a New Database in Aura

Login - or create a free account

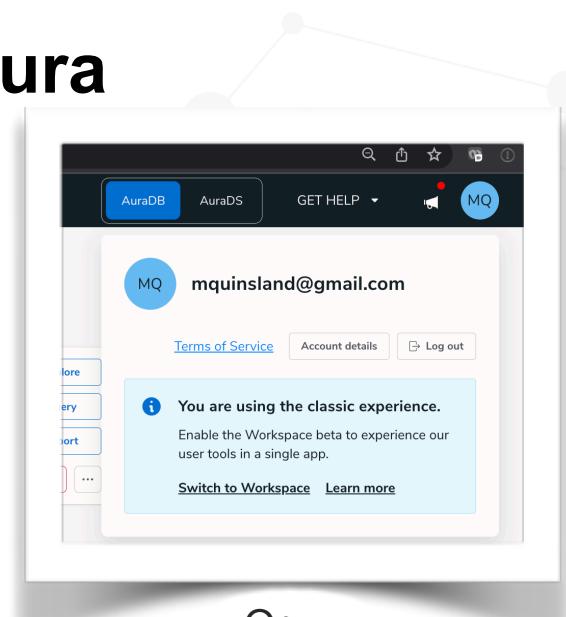
- Proceed to <https://login.neo4j.com>
- Enter username and password
- Create a new account if necessary



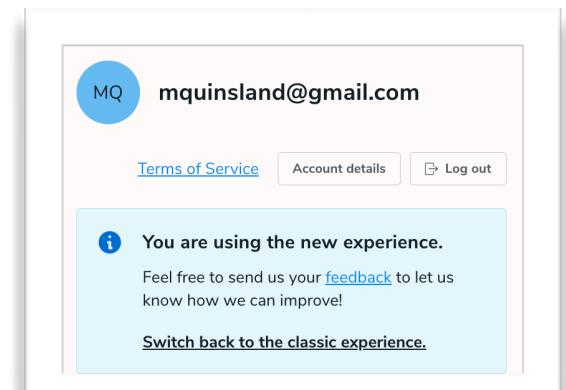
Lab 0.1 - Creating a New Database in Aura

Switch to New Workspace Experience

- After logging in, click on your initials in the menu bar
- You will see one of these messages to indicate whether you are in the classic version, or the new Workspace version.
- We will be using the new Workspace version. (it's in beta testing so things may get rough!)



Or



Lab 0.2 - Creating a New Database in Aura

Create a new Instance

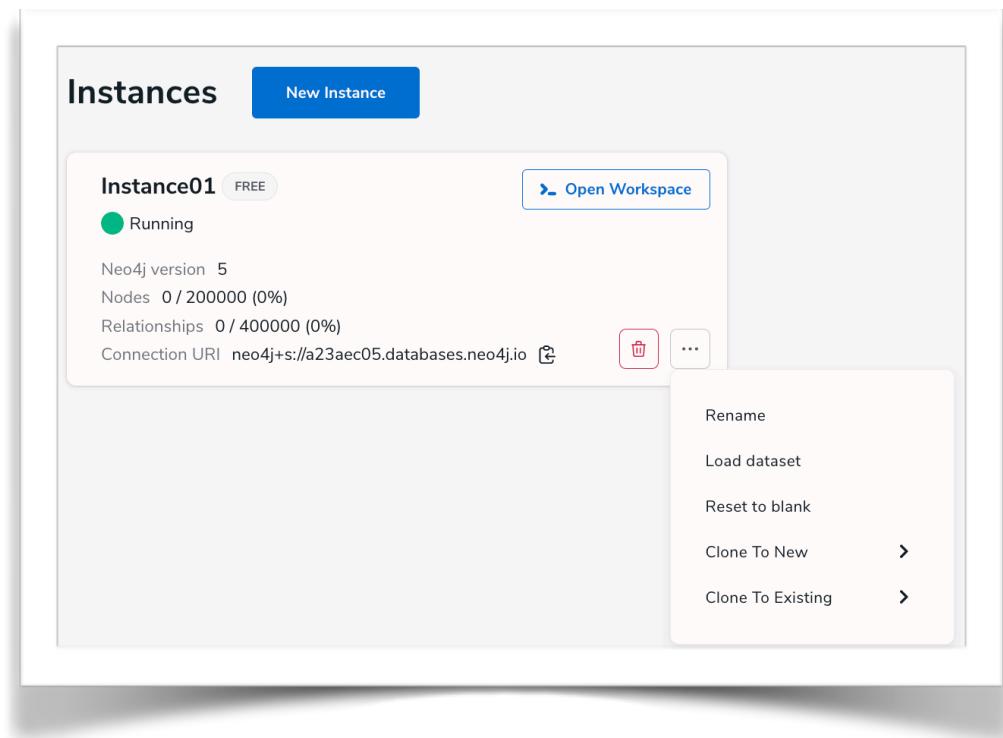
- Click on the New Instance Button
- Select the Empty instance
- Your new instance will be named Instance01 and a password will be generated. Copy the password to the clipboard and download the password file
- Please remember where you downloaded the password file!

The screenshot shows two overlapping interface panels. The top panel is titled 'Instances' and features a prominent blue 'New Instance' button. The bottom panel is titled 'Credentials for Instance01' and contains fields for 'Username' (set to 'neo4j') and a 'Generated password' (a long string of characters). A 'Download' button is located next to the password field. Below these fields is a light blue box containing the text: 'We strongly advise changing this initial password.' At the bottom of the panel is a checkbox with the text: 'I confirm I have copied or downloaded the above credentials, as this password will not be available after this point.' A 'Continue' button is located at the very bottom right of the panel.

Lab 0.3 - Creating a New Database in Aura

Rename the new Instance

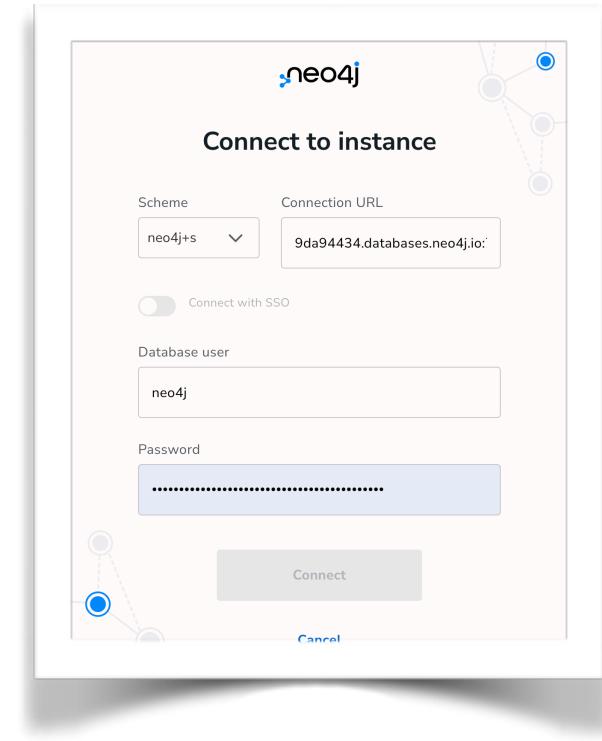
- In a moment or two, the instance will be created and the status will change to 'Running'
- Change the name to something useful by clicking on the ... button next to the trash can.
- Some basic statistics will be shown
- The instance will be free to use for its lifetime. After periods of inactivity it will be suspended, but resuming it requires only a simple click.



Lab 0.4 - Creating a New Database in Aura

Testing the Connection

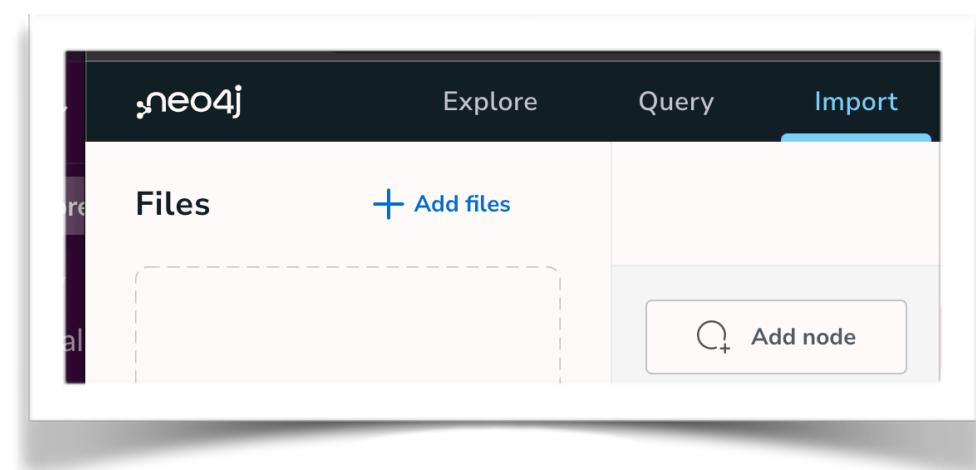
- Click on the Open Workspace button to display the Connect to Instance dialog.
- Paste your password into the appropriate field.
- If you forgot your password and failed to download it, you must delete the instance and start over.



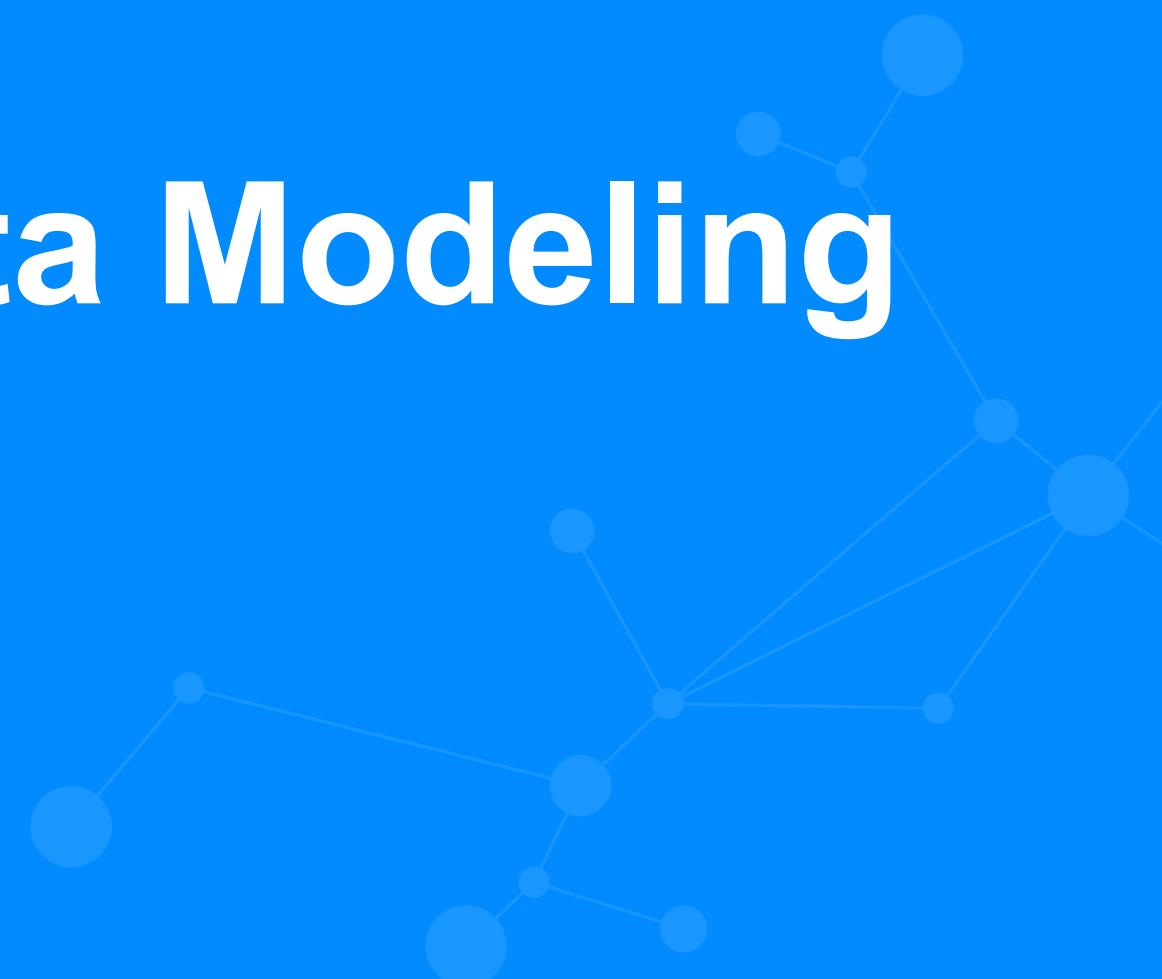
Lab 0.5 - Creating a New Database in Aura

Viewing the Workspace

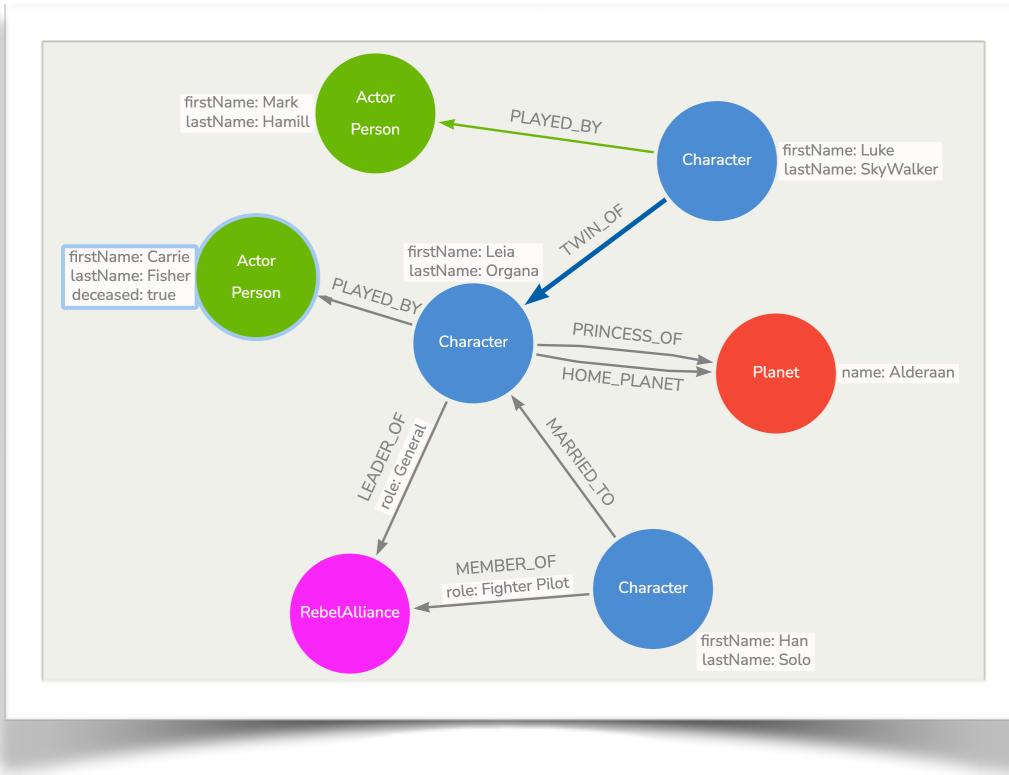
- The Aura Developer Workspace has 3 main tabs.
- Explore: Uses Neo4j Bloom to analyze and visualize the graph without writing code
- Query: Write and execute Cypher queries to create, read, update, and delete data.
- Import: Use the Neo4j Import tool to load CSV files without writing code
- We will start on the Query tab
- But first, let's create some data and a data model!



Neo4j Data Modeling



Visually Create Data Models with Arrows

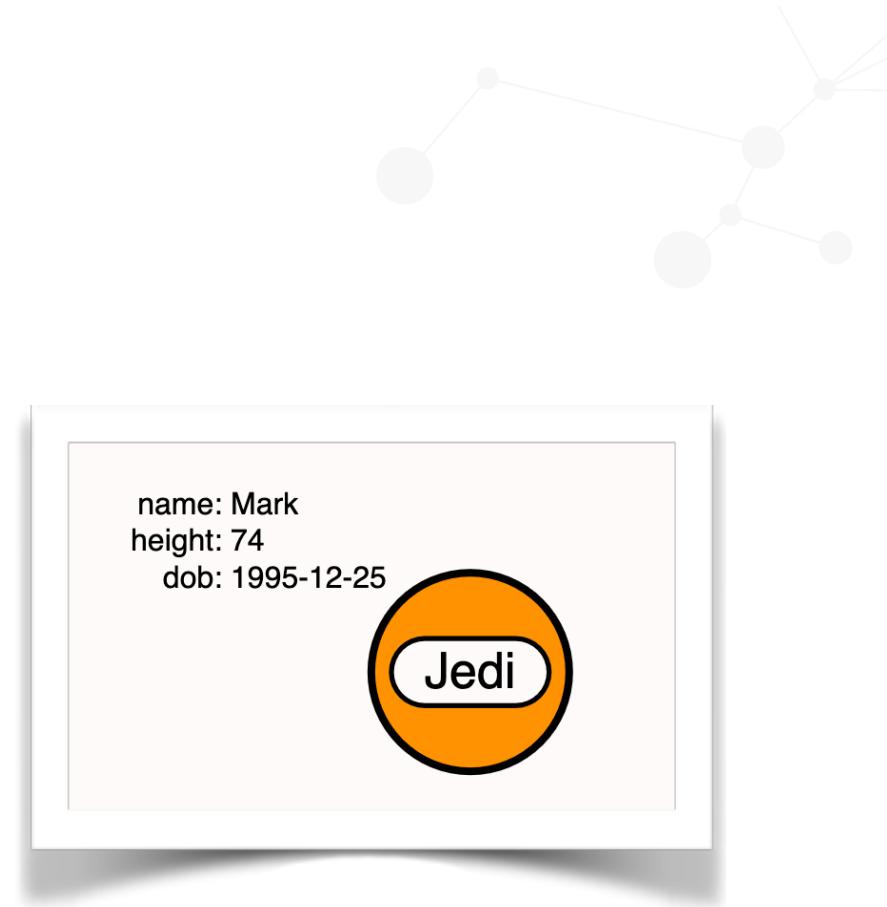


- Create Graphs visually using sample data
- Export to Cypher
- Import / Export models
- Simple to learn
- <https://arrows.app>

Lab 1 - Data Modeling with Arrows

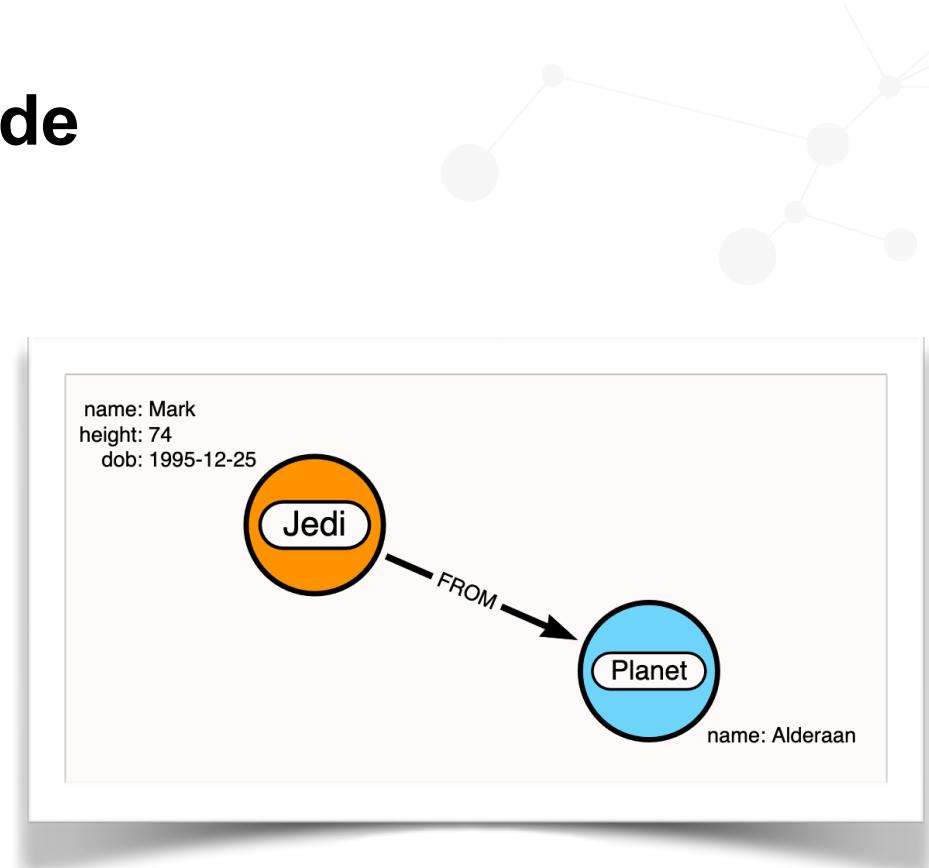
Lab 1.1 - Creating a New Node

- Bring up <https://arrows.app> on a browser
- Tap on the empty node to highlight it
- On the properties tab, give the new node a label (not a caption)
- Give the node a property:value pair (name: your name)
- Give the node 2-3 additional properties
- Give the node a color



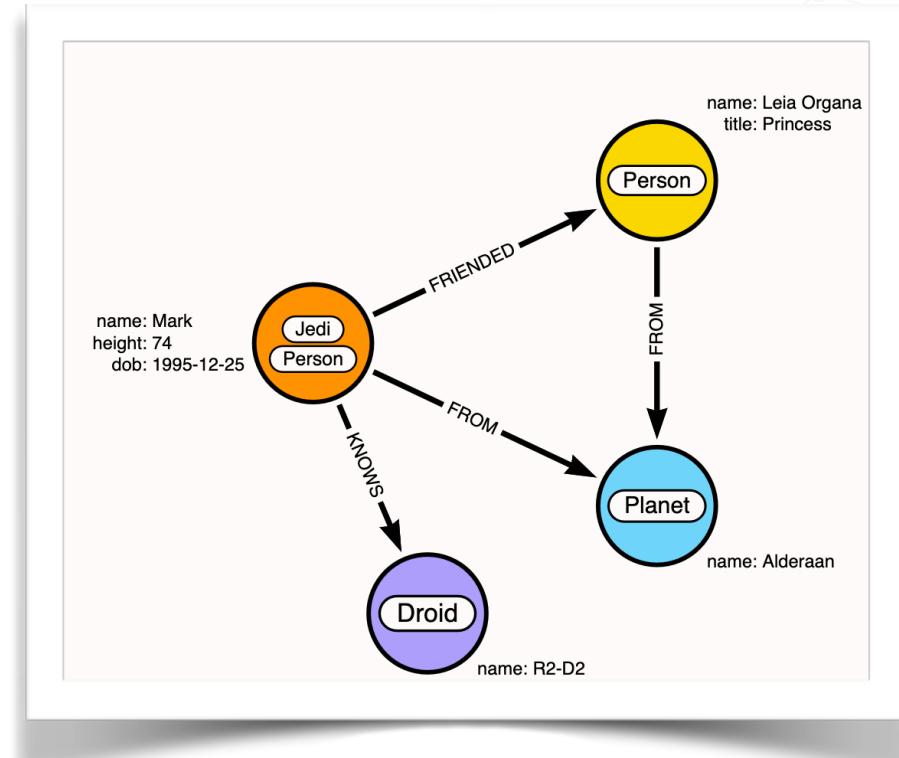
Lab 1.2 - Creating a Related Node

1. Hover over the border of the source node until the border shows that it is selected
2. Click on the selection border and drag it away from the node. You will see a new node appear.
3. Drop the new node away from other nodes
4. On the properties tab, give the new node a label (not a caption)
5. Give the node a property:value pair (name: your name)
6. Select the relationship and give it a type.



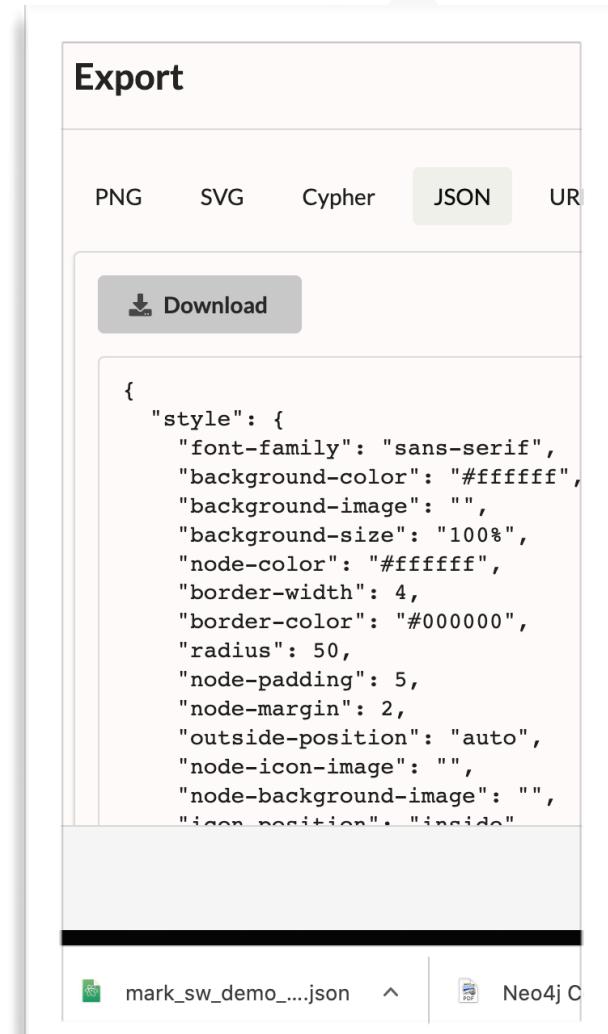
Lab 1.3 - Creating Additional Related Nodes

1. Hover over the border of the source node until the border shows that it is selected
2. Click on the selection border and drag it away from the node. You will see a new node appear.
3. Drop the new node away from other nodes
4. On the properties tab, give the new node a label (not a caption)
5. Give the node a property:value pair (name: your name)
6. Select the relationship and give it a type.
7. Repeat for 2-3 more nodes



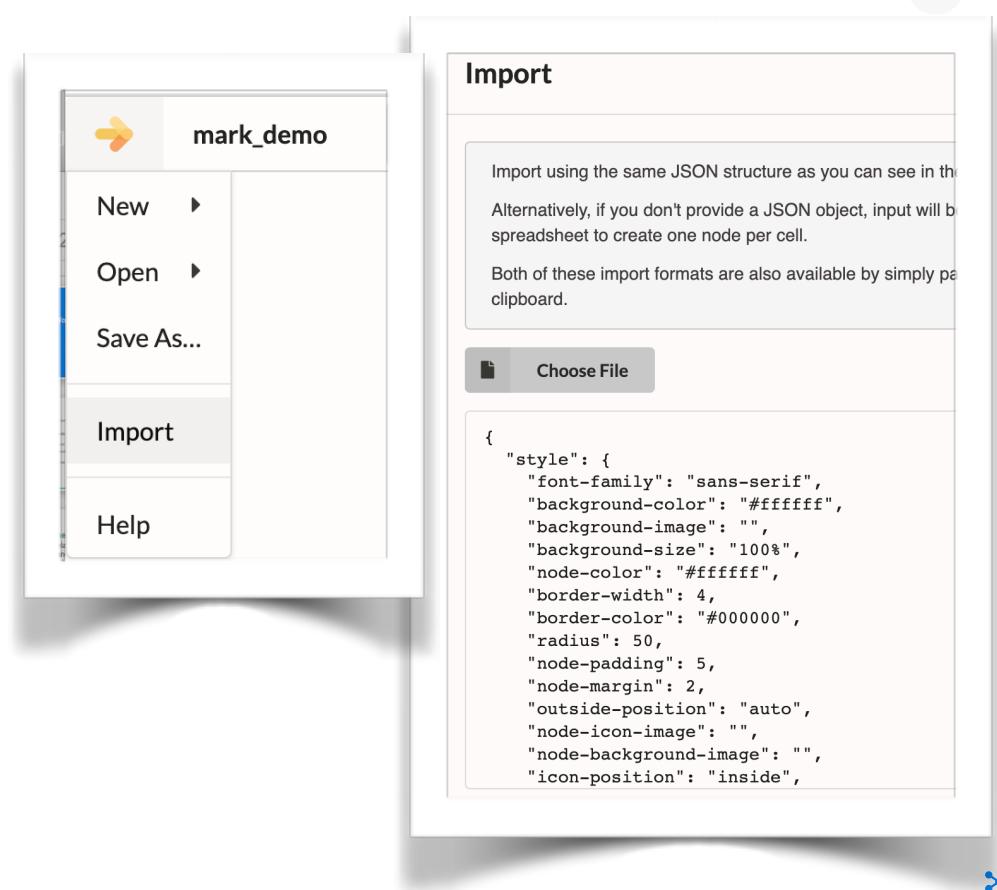
Lab 1.4 - Exporting a Model to JSON

1. Give the model a name
2. Click on the Download/Export button
3. Select the JSON tab
4. Click on the Download button
5. Select a safe place to store the JSON file.



Lab 1.5 - Import a Previously Exported Model

1. Click on the yellow arrow at the top left corner
2. Select the Import item
3. Choose the JSON file to be uploaded. Use the json file "starwars_intro_model.json" from the downloaded zip file.
4. View the new nodes / relationships.



Lab 1.6 - Exporting A Model to Cypher

1. Click on the Download/Export button
2. Select the Cypher tab
3. Select the CREATE option
4. Review the Cypher code
5. Select the MERGE option
6. Compare this Cypher with the previous



The screenshot shows the Neo4j browser interface with an "Export" dialog open. The dialog has tabs for PNG, SVG, Cypher (which is selected), JSON, URL, and GraphQL. Below the tabs, there's a section for "Cypher Clause" with radio buttons for CREATE, MATCH, and MERGE, where MERGE is selected. A note below says: "MERGE query behaviour depends on what data is already present in the database. You may need to edit the query to achieve exactly the behaviour you are looking for. Please see MERGE documentation for guidance." At the bottom are "Copy to clipboard" and "Run in Neo4j Browser" buttons. The Cypher code shown is:

```
MERGE (n2:Person {name: "Leia Organa", title: "Princess"})<-[ :FRIENDED ]-(n0:Jedi:Person {name: "Mark", height: 74, dob: "1995-12-25"})-[:FROM]->(:Planet {name: "Alderaan"})-[:FROM]-(n2)
MERGE (n0)-[:KNOWS]->(:Droid {name: "R2-D2"})
```

Lab 1.7 - Importing the data into Neo4j using the browser

1. Copy the Cypher text that was generated by the Export window
2. Switch to your browser and bring up the Aura Developer Workbench tab
3. Paste the code into the execution window
4. Click on the small blue arrow to execute the command

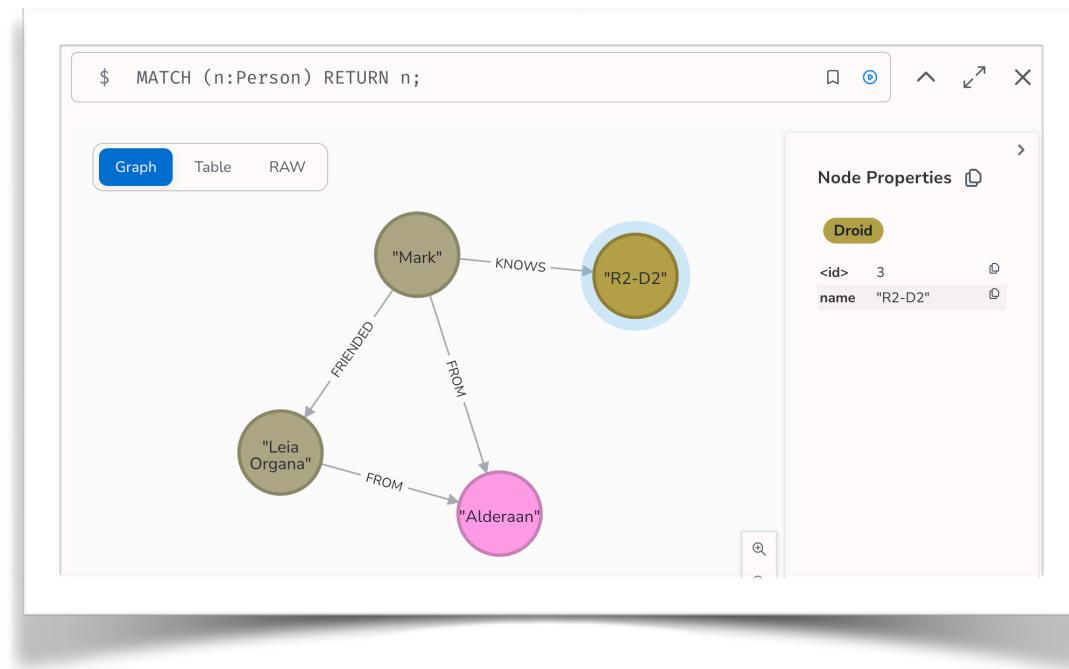
```
1 CREATE (n6:Person {name: "Leia Organa", title:  
"Princess"})-[:FRIENDED]-(n4:Jedi:Person {name:  
"Mark", height: 74, dob: "1995-12-25"})-[:FROM]->  
(:Planet {name: "Alderaan"})-[:FROM]-(n6),  
2 (n4)-[:KNOWS]->(:Droid {name: "R2-D2"})
```



Lab 1.8 - Viewing the data in a Neo4j Browser

1. Run a simple Cypher command to view your data
2. Double-click on the nodes to see related nodes appear.

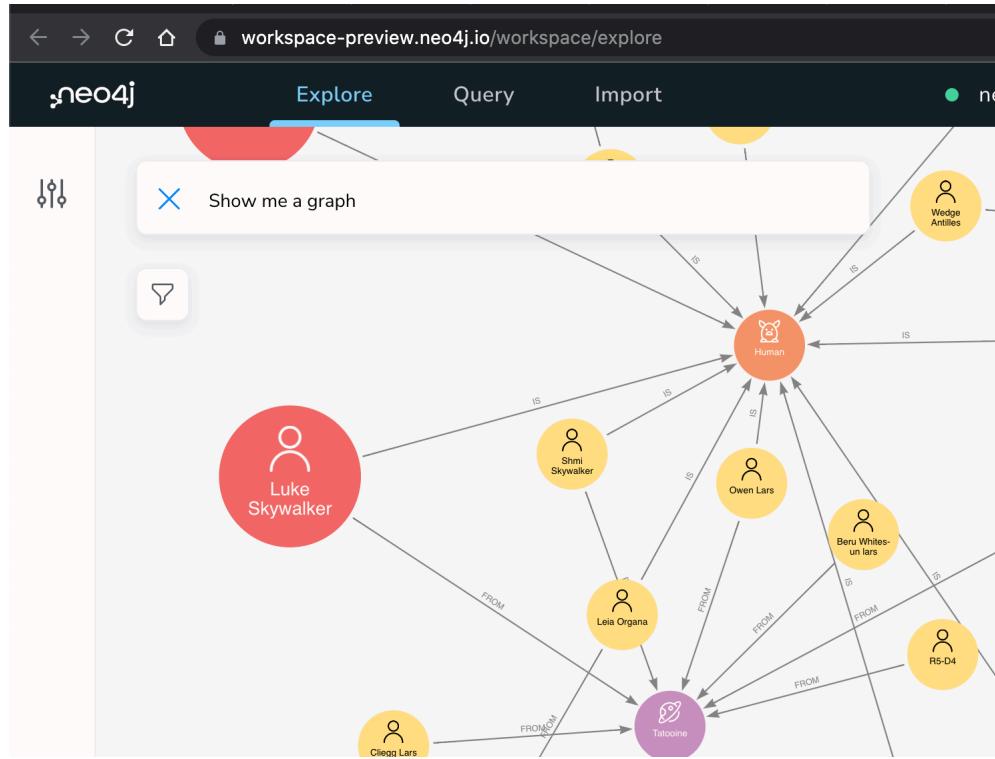
```
MATCH (n:Person) RETURN n;
```



Neo4j Workspace



Explore, Query, Import with Neo4j Workspace



- Easy-to-use work surface (beta)
- Import CSV files
- Edit/Run Cypher Commands
- Visually explore the graph without using Cypher

Run Cypher Commands -Workspace - Query

neo4j Explore Query Import neo4j+s://d44b75fb.databases.neo4j.io:7687 Send feedback ?

Database Information

Nodes (255) Character Planet Species User

Relationships (657) FROM HOMeworld INTERACTS_WITH IS

Property keys average_height average_lifespan betweenness birthYear classification climate closeness count data designation diameter eye_colors eyecolor gender gravity hair_colors hairColor height homePlanet homeworld

Show all (19 more)

1 MATCH (c3po:Character {name:"C-3PO"}),
(r2:Character {name:"R2-D2"})
2 MATCH p=(c3po)-[i:INTERACTS_WITH]-(others)
3 WHERE NOT (others)-[:INTERACTS_WITH]-(r2)
4 return p

Graph Table RAW

Results overview

Nodes (15) Character (15)

Relationship (14) INTERACTS_WITH (14)

Started streaming 14 records after 3ms and completed after 13ms.

Import CSV files -Workspace - Importer

Preview Run import ...

Add node Delete

Mapping details

Type HAS_CHARACTER Switch direction

File sw_nodes_credits.csv ▼

Filter file

Node	ID property	ID file column
From	Credit	creditId creditId ▼
To	Character	charact... charact... ▼

Need help?

Definition File mapping File mapping

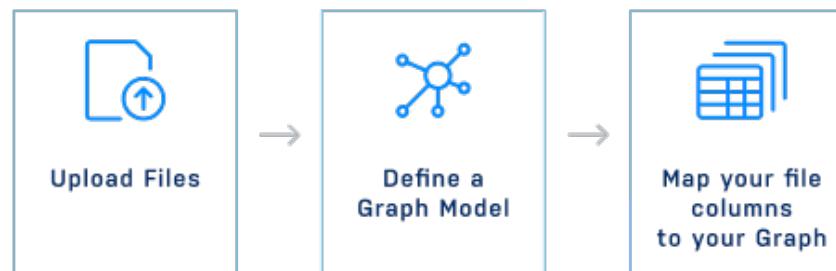
+ Select from file + Add Property



Neo4j Data Importer

Small no-code UI for loading flat file data

- Beta on AuraDB
- Not designed to replace ETL tools or production imports
- Support for small flat-file inputs (CSV, TSV)
- The ability to sketch out a graph model and map your input data
- Loads into any Neo4j database reachable from your machine
- Data is only kept in your web browser; the tool doesn't use a server-side component
- Self-hosted may use this website - <https://data-importer.graphapp.io/>

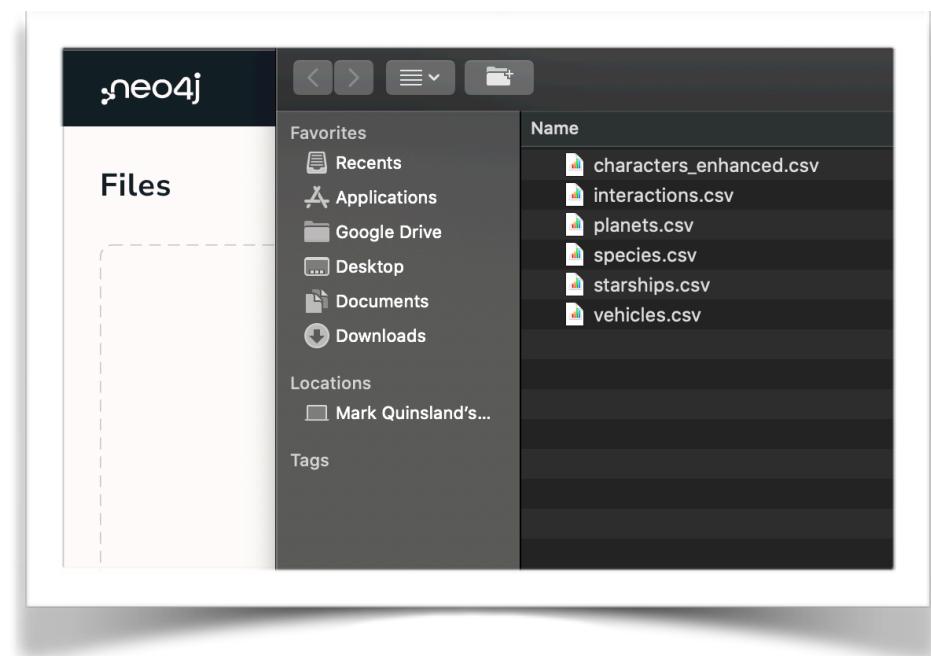


Lab 2 - Using Workspace Importer

Lab 2.1 - Using Workspace Importer

Selecting CSV files to Import

- Download the data files from
- https://github.com/mquinz/star_wars_demo/blob/master/data/star_wars_data.zip
- Unzip the files
- Click on the + Add files link on the Workspace Import tab or drag-and-drop the files onto the files tab.



Lab 2.2 - Previewing Data from CSV Files

- After adding the 6 CSV files, expand/collapse the file names to preview the data in the files
- The importer will attempt to determine data types for the values, but you can override/correct this in the next step.

^ characters_enhanced.csv	...
^ interactions.csv	...
^ planets.csv	...
▼ species.csv	...
name	Hutt
classification	gastropod
designation	sentient
average_height	300
skin_colors	green, brown, tan
hair_colors	NA
eye_colors	yellow, red
average_lifespan	1000
language	Huttese
homeworld	Nal Hutta
^ starships.csv	...
^ vehicles.csv	...

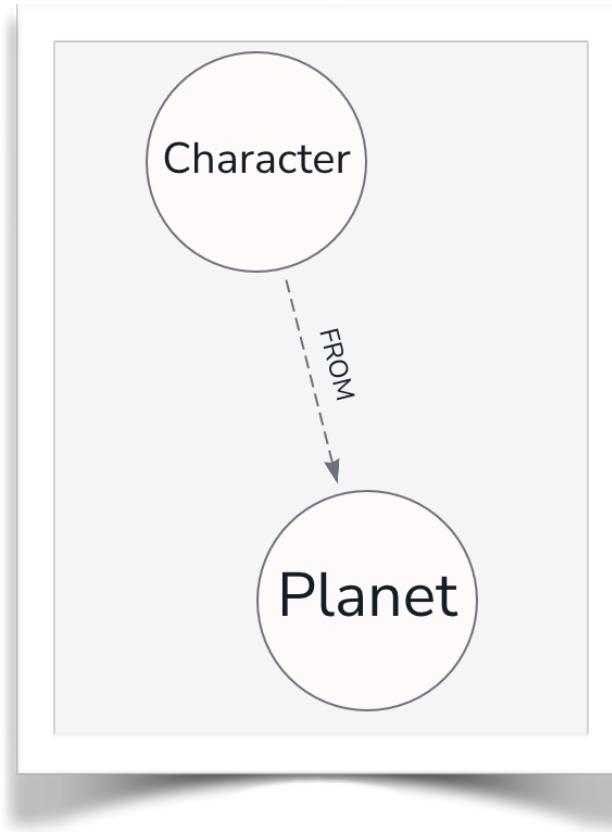
Lab 2.3 - Adding Nodes from CSV File

- Click on the Add Node Button and an empty node type will appear.
- In the Mapping Details tab, enter the name of a Label e.g. Character
- Select a CSV file that contains the data to be loaded for that Label
- Click on the ‘Select from file’ link and choose the property names to be loaded.
- Edit data types or field names
- In the ID field, choose which property value should be used as the primary key of each node

The screenshot shows the Neo4j Import dialog with the 'Mapping Details' tab selected. On the left, a tree view lists CSV files: 'characters_enhanced.csv' (expanded), 'interactions.csv' (collapsed), and 'planets.csv' (collapsed). Under 'characters_enhanced.csv', properties like name, gender, birthYear, etc., are listed with their values. A 'Select from file' button is shown, with a 'Select all' checkbox checked. On the right, the 'Label' field is set to 'Character'. Below it, the 'File' dropdown shows 'characters_enhanced.csv'. The 'Properties' tab is active, displaying the properties and their types (e.g., name: string, gender: string) with edit and delete icons. The 'Mapping' tab is also visible. At the bottom, there are 'Cancel' and 'Confirm' buttons.

Lab 2.4 - Creating Relationships Between Nodes

- Select a Node type that will serve as the 'From' node
- Hover over its border until the border color changes and a + indicator appears.
- Drag the + indicator to the node that will be the 'To node.'
- Double-Click on the relationship line to select it and type in the relationship type
- The line will appear broken until the relationship info is completed



Lab 2.5 - Configuring New Relationships

- Complete the remaining fields on the Mapping Details tab.
- Type: The name of the relationship
- File: Which file has the source/target information for the relationship
- From: Which field contains the ID of the From node
- To: Which field contains the ID of the To node
- Properties: Any fields that should be used as relationship properties.

Mapping Details

Type *i* ↵
FROM

File *i* ↴
characters_enhanced.csv

From *i* Character (name)
name

To *i* Planet (name)
homePlanet

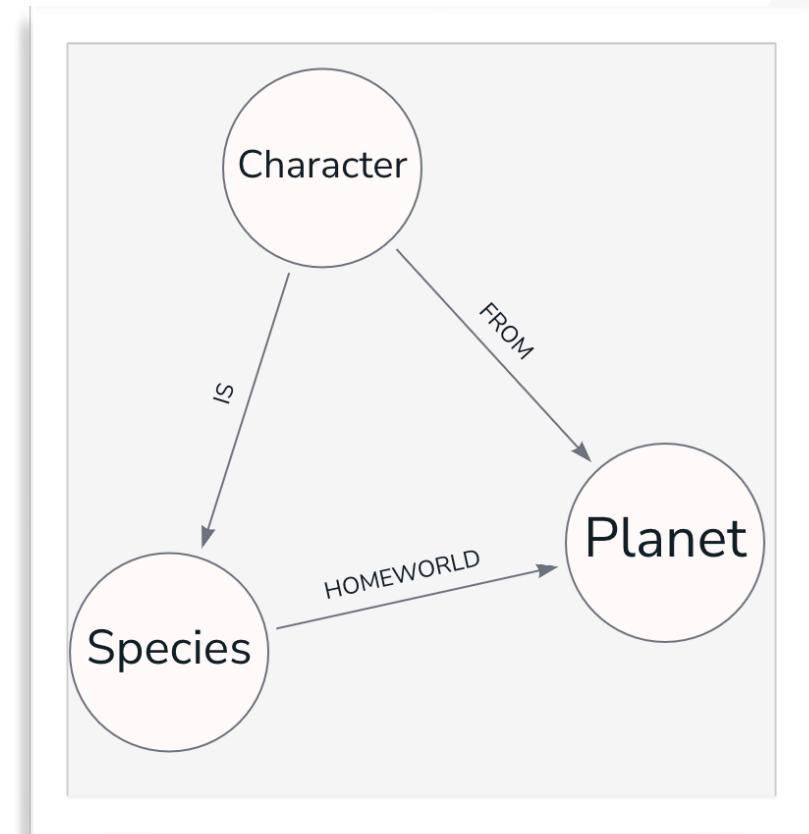
Properties Mapping

+ Select from file + Add new

No properties defined

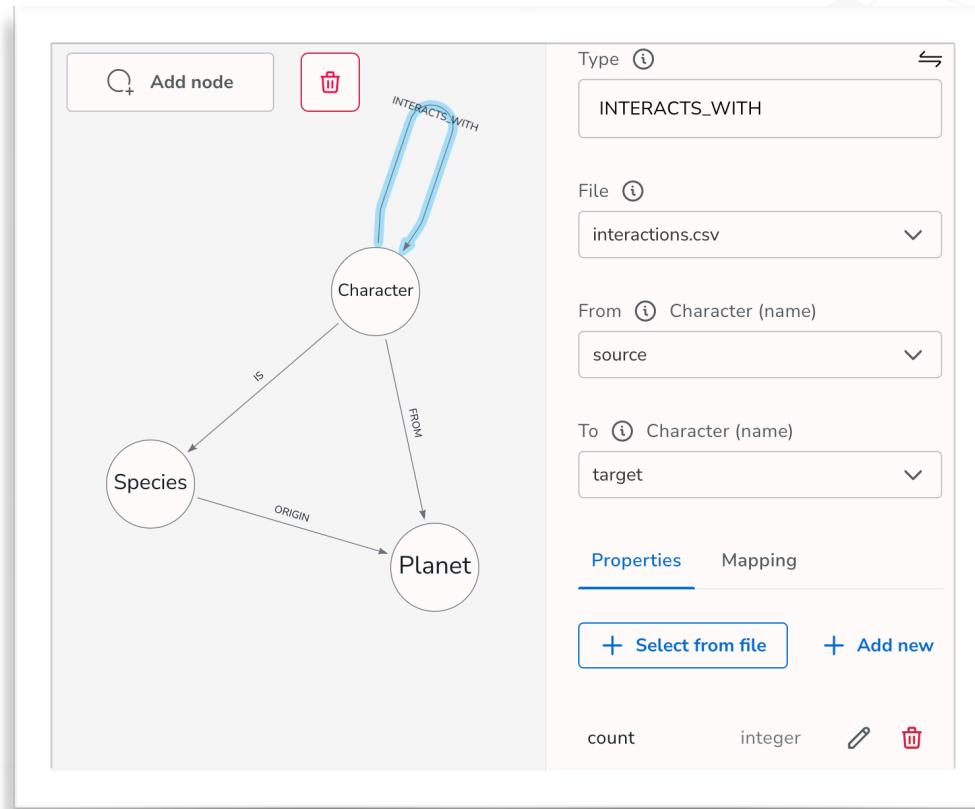
Lab 2.6 - Load Species Nodes and Relationships

- Load species.csv file
- Use species name as ID field
- Create relationship from Character to Species.
- Create relationship from Character to Planet



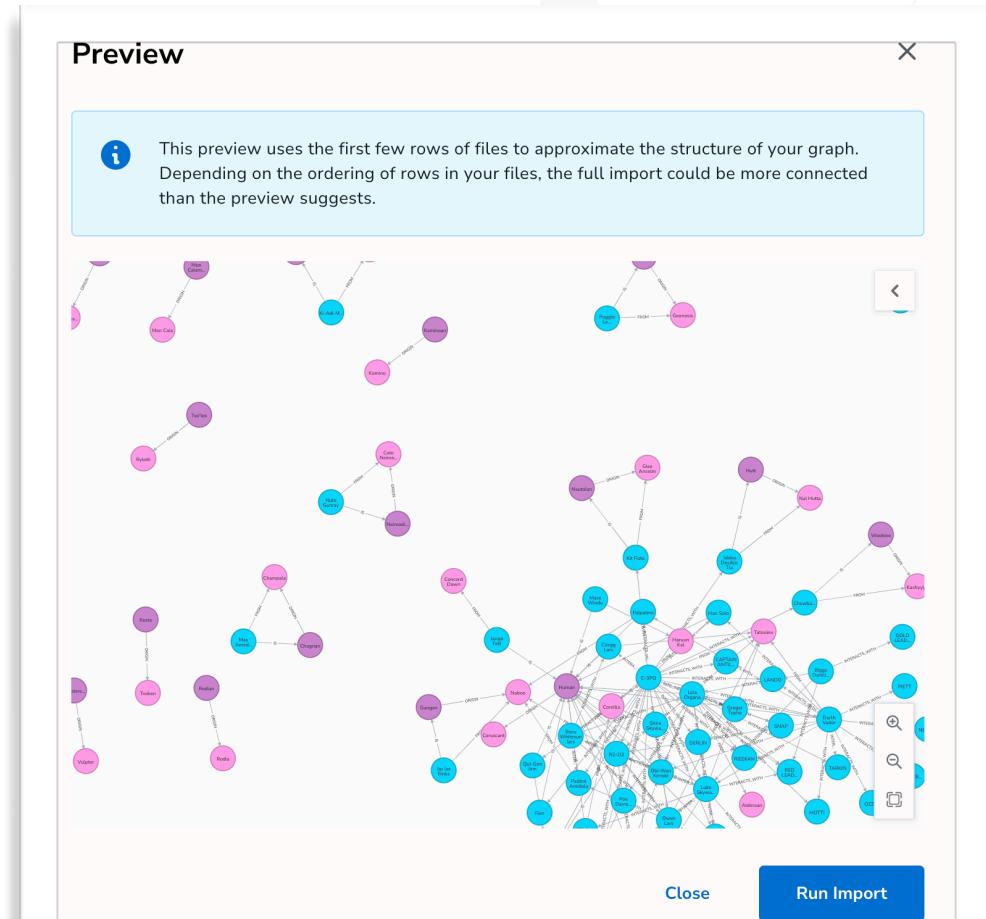
Lab 2.7 - Creating Self-Join Relationships for Characters

- Select the Character node and hover over its border until the + indicator appears
- Drag the selection + indicator away from the Character node, and then drag it back and drop it on the Character node.
- Fill out the mapping details.
- Type: INTERACTS_WITH
- File: interactions.csv
- From: source
- To: target
- Properties: count



Lab 2.8 - Preview the Import

- Click on the Preview button
- A sampling of records from each of the files will used to give an idea of how the nodes and relationships will look.
- If everything looks okay, click on the Run Import button.



Lab 2.9 - Run the Import and Review Results

- Click on the Run Import button
- A summary of the import will be shown.
- Each Node Label and each Relationship Type will be summarized.
- If all looks good, click on the Start Exploring button.

Import results

Import completed without errors • Total time: 00:00:04

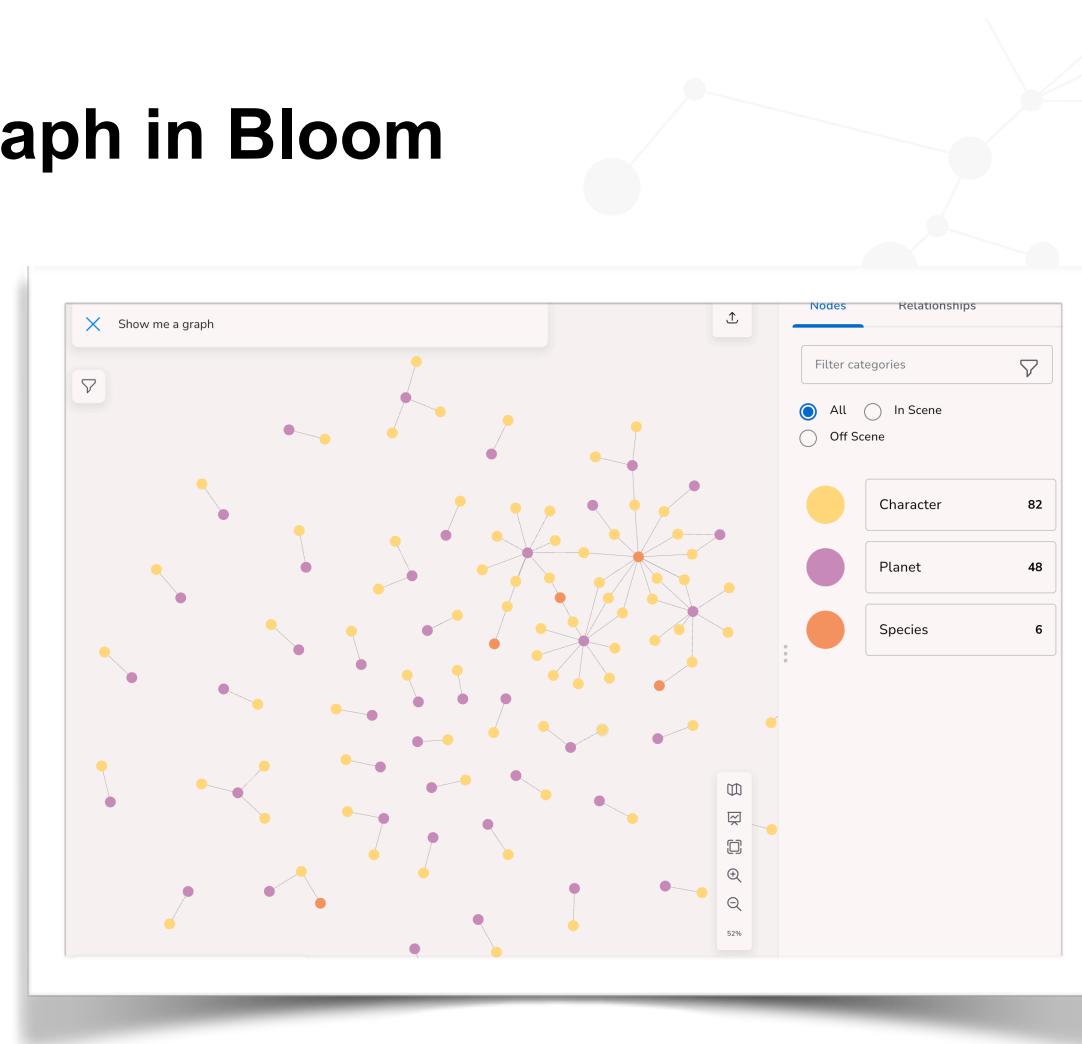
Species	species.csv	Show cypher					
Time Taken	File Size	File Rows	Nodes Created	Properties Set	Labels Added	Query Count	Query Time
00:00:01	3.2 KiB	37	37	346	37	1	00:00:04

FROM	characters_enhanced.csv	Show query				
Time Taken	File Size	File Rows	Relationships Created	Properties Set	Query Count	Query Time
00:00:00	16.0 KiB	151	82	0	1	00:00:00

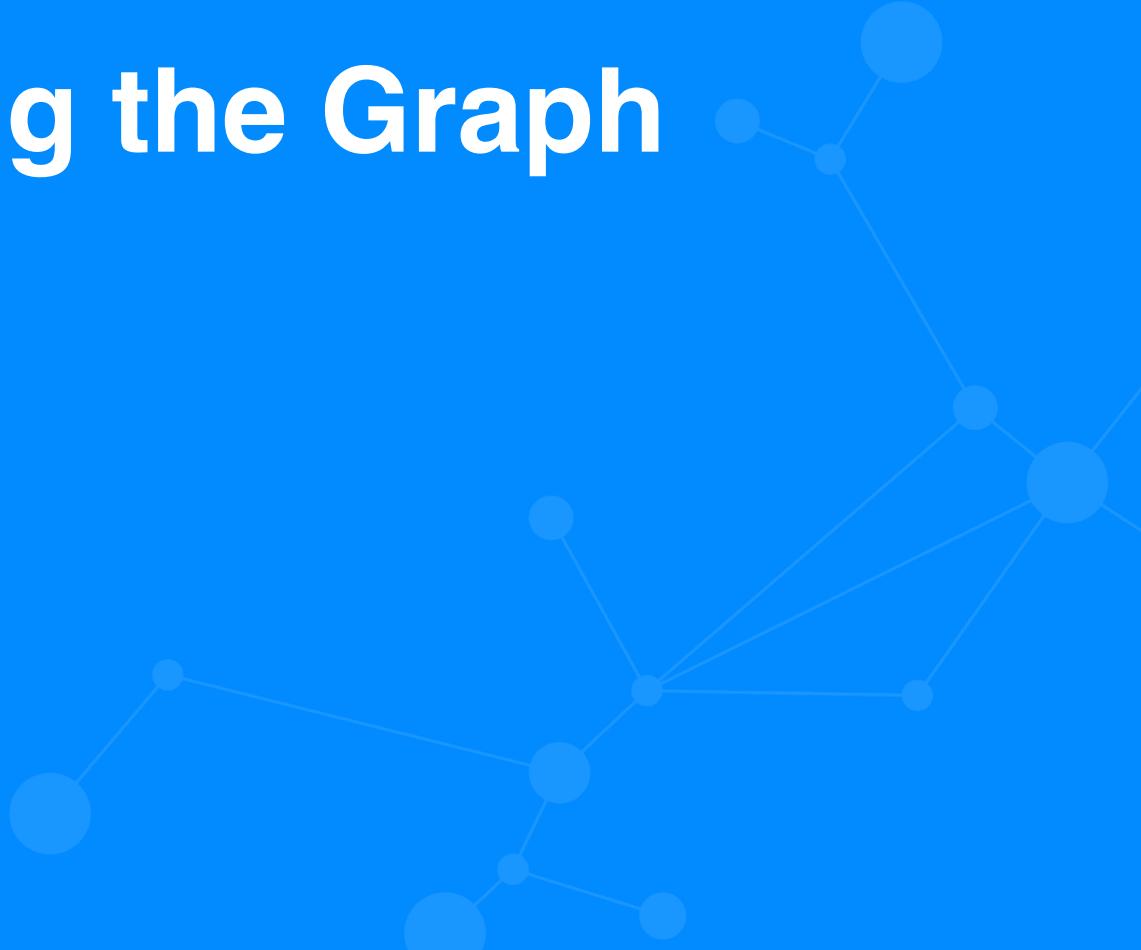
[Close](#) [Start Exploring](#)

Lab 2.9.1 - Explore your Graph in Bloom

- Click on the Start Exploring button
- This will show the Explore tab which uses Neo4j Bloom to display the newly imported data



Lab 3 - Querying the Graph



Lab 3.1 - Basic Queries - The Match Statement

- Most queries start with a MATCH statement which is used for pattern matching.
- It is generally the most important part of the query
- All nodes/relationships/paths that meet the filtering criteria are selected.
- Selected items are then passed to the next part of the query

```
// Return all characters
MATCH (c:Character)
RETURN c
```

```
// Return a specific character
MATCH (c:Character)
WHERE c.name = "Han Solo"
RETURN c
```

Lab 3.2 - Basic Queries - The Match Statement

- Most queries start with a MATCH statement
- Used for pattern matching - similar in some ways to a WHERE clause
- All nodes/relationships/paths that meet the filtering criteria are selected
- Selected items are then passed to the next part of the query
- It is generally the most important part of the query

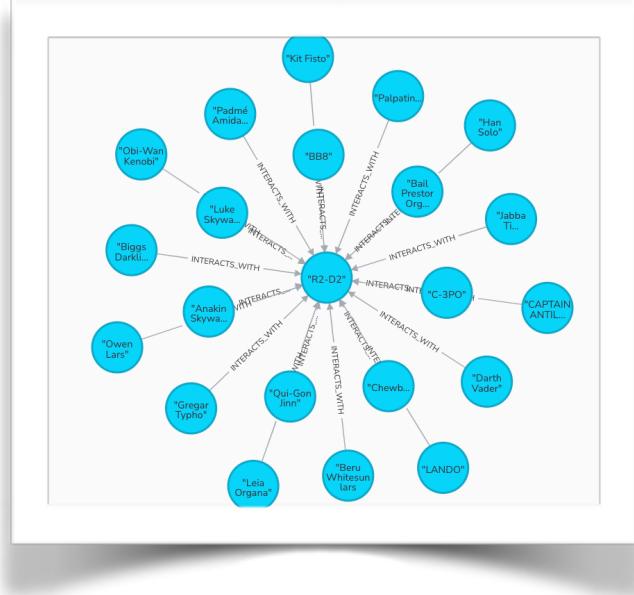
```
// Return all Characters from Alderaan
MATCH (c:Character)-[:FROM]->(p:Planet)
WHERE p.name = "Alderaan"
RETURN c
```

```
// Return a specific Character
MATCH (c:Character)
WHERE c.name = "Han Solo"
RETURN c
```

Lab 3.4 - Basic Queries - The Match Statement

Node Properties Used as a Filter

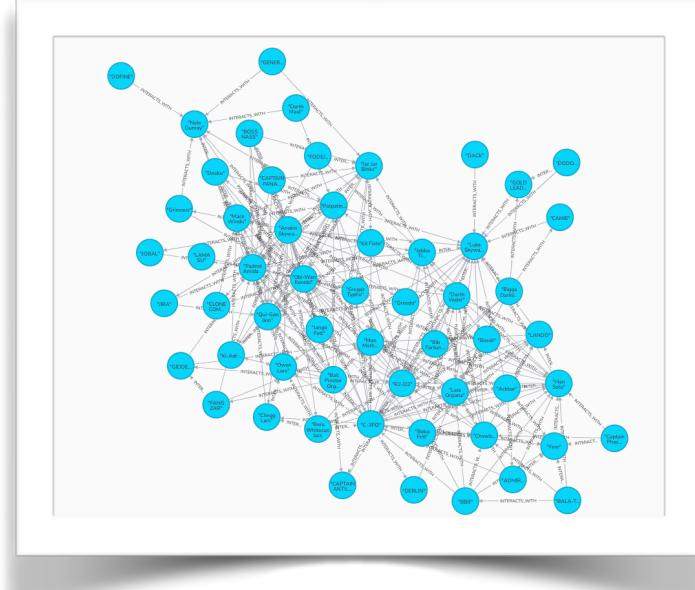
```
// Return all Characters that interacted with R2-D2
MATCH (c:Character)-[INTERACTS_WITH]-(other:Character)
WHERE c.name = "R2-D2"
RETURN other
```



Lab 3.5- Basic Queries - The Match Statement

Unlike SQL, Variable-Length Queries are simple

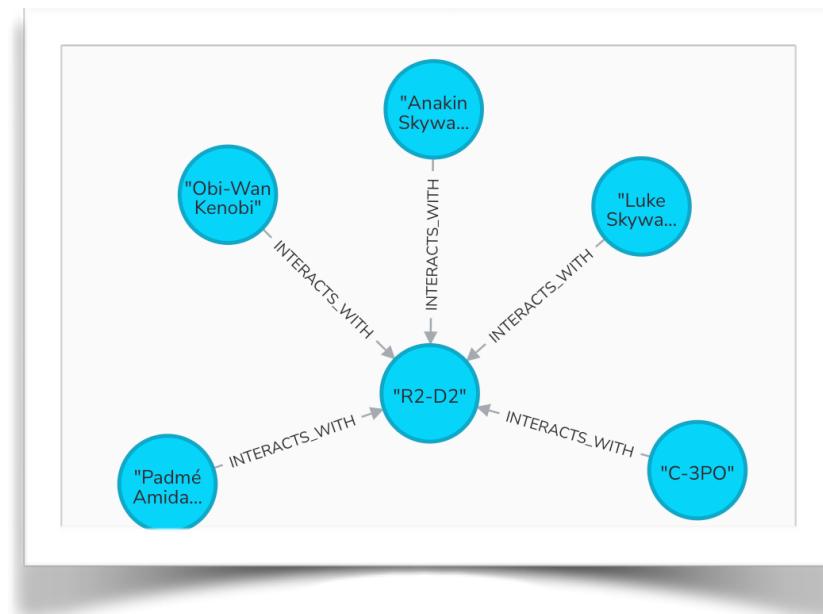
```
// Variable length query - interacts_with 3 levels down
MATCH p=(c:Character)->[INTERACTS_WITH*..3]-(other:Character)
WHERE c.name = "R2-D2"
RETURN p
```



Lab 3.6 - Basic Queries - The Match Statement

Relationship Properties Can Also be Used

```
// Who does R2-D2 interact with the most
MATCH p=(c:Character)-[i:INTERACTS_WITH]-(other:Character)
WHERE c.name = "R2-D2"
RETURN p ORDER BY i.count DESC LIMIT 5
```



Lab 3.7 - Basic Queries - The Match Statement

Aggregation Values Used as a Filter

```
// Which characters interact with the most characters
MATCH (c:Character)-[INTERACTS_WITH]-(other)
WITH c.name as character, count(other) as others
WHERE others > 20
RETURN character, others
ORDER BY others desc LIMIT 5
```

character	others
"Anakin Skywalker"	44
"Obi-Wan Kenobi"	39
"C-3PO"	38
"Padm� Amidala"	36

Showing 1-5 of 5 results

Lab 3.8 - Basic Queries - The Match Statement

String Comparison Used as a Filter

```
// Which characters have Darth as part of their name
MATCH (c:Character)
WHERE c.name STARTS WITH "Darth"
RETURN c
```



Lab 3.8 - Basic Queries - The Match Statement

String Comparison Used as a Filter

```
// Which characters have Darth as part of their name
MATCH (c:Character)
WHERE c.name STARTS WITH "Darth"
RETURN c
```



Lab 4 - Updating the Graph



Lab 4.1 - Update Queries - Creating New Nodes

Using the CREATE Statement

- New nodes can be created using the CREATE statement
- Unless uniqueness constraints are used, duplicates may result

```
// Create a new Character  
CREATE (c:Character {name:"Mark Q",  
height:74,hair_color:"Fleshtone"})  
RETURN c
```



"Mark Q"

Lab 4.2 - Update Queries - Creating New Nodes

Using the MERGE Statement

- New nodes can also be created using the MERGE statement
- Neo will check first to see if any nodes match before creating a new node

```
// Carefully create a new Character using MERGE
MERGE (c:Character {name:"Mark Q"})
ON CREATE SET c +=
{height:74,hair_colors:"Fleshtone"}
RETURN c
```



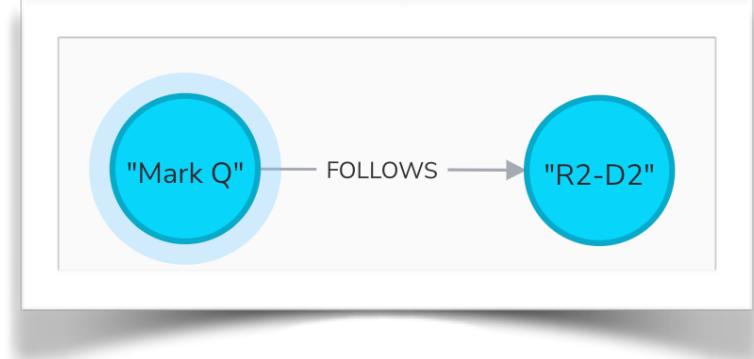
"Mark Q"

Lab 4.3 - Update Queries - Creating New Relationships

Using the MERGE Statement

- New relationships can also be created using the MERGE statement
- Use the MATCH statement to find the “from” and “to” nodes
- Use the MERGE statement to create a new relationship (if necessary)

```
MATCH (r2:Character), (me:Character)
WHERE r2.name = 'R2-D2' AND me.name = 'Mark Q'
MERGE (me)-[r:FOLLOWS]->(r2)
RETURN r2,r,me
```



Exercise 4.3.1 - Creating a New Node

Using the MERGE Statement

- Create a New node with a name of your choosing using a MERGE statement.
- Add another property by adding a key/value pairs to the map
- Run the statement again to ensure that the record

```
// Carefully create a new Character using MERGE
MERGE (c:Character {name:"Mark Q"})
ON CREATE SET c += {height:74,
    hair_colors:"Fleshtone",
    jedi:true }
ON MATCH SET c.updated = date()
RETURN c
```



Exercise 4.3.2 - Creating a New Relationship

Using the MERGE Statement

- Create a relationship from your new node to an existing Film

```
// Add your new Jedi to a Film
MATCH (c:Character {name:"Mark Q"}),
MATCH (f:Film {name:"The Empire Strikes Back"})
MERGE (c)<-[r:HAS_CHARACTER]-(f)
RETURN c, r, f
```

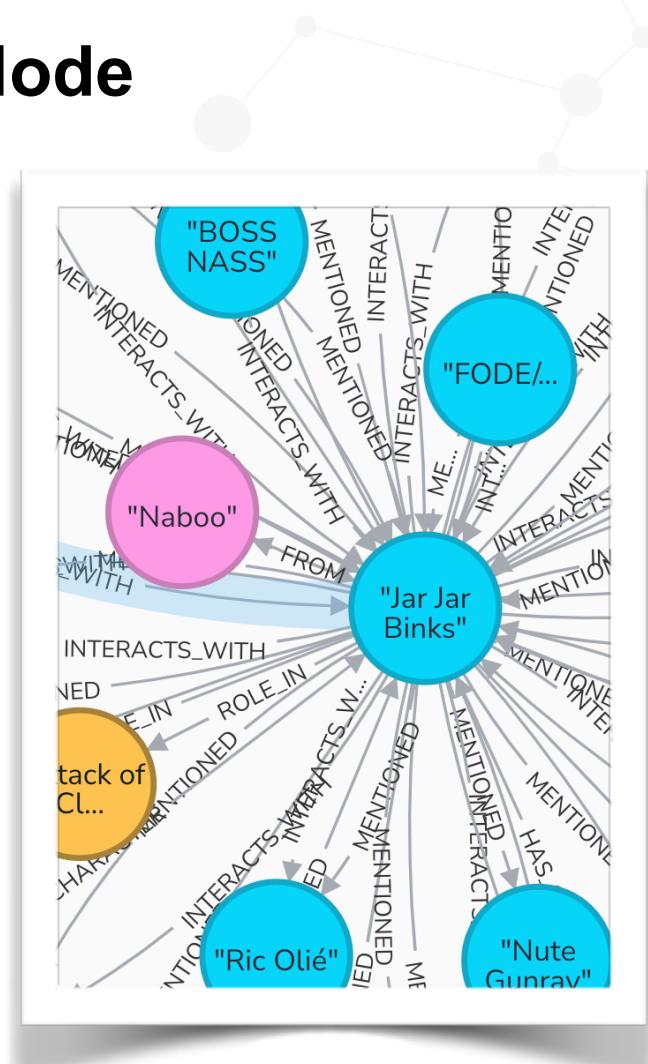


Lab 4.5 - Update Queries - Deleting a Node

Using the DELETE Statement

- Nodes may be deleted only if they do not have any relationships.
- DETACH DELETE will automatically remove any relationships so the node may be deleted.
- USE EXTREME CAUTION

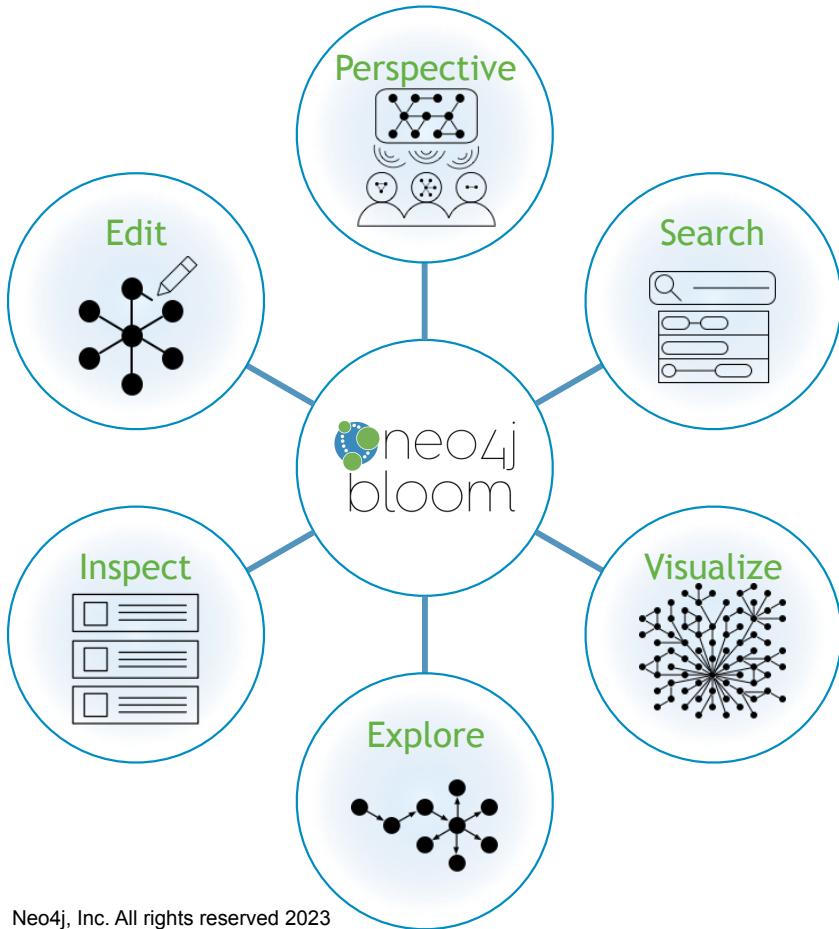
```
// Find the Node to be deleted
MATCH (c:Character WHERE c.name = 'Jar Jar Binks')
// Delete the object and all of its relationships
DETACH DELETE c
```



Neo4j Bloom



Explore & Collaborate with Neo4j Bloom



Explore Graphs Visually

Prototype Concepts Faster

Collaborate Across Teams

Neo4j Bloom's Intuitive User Interface

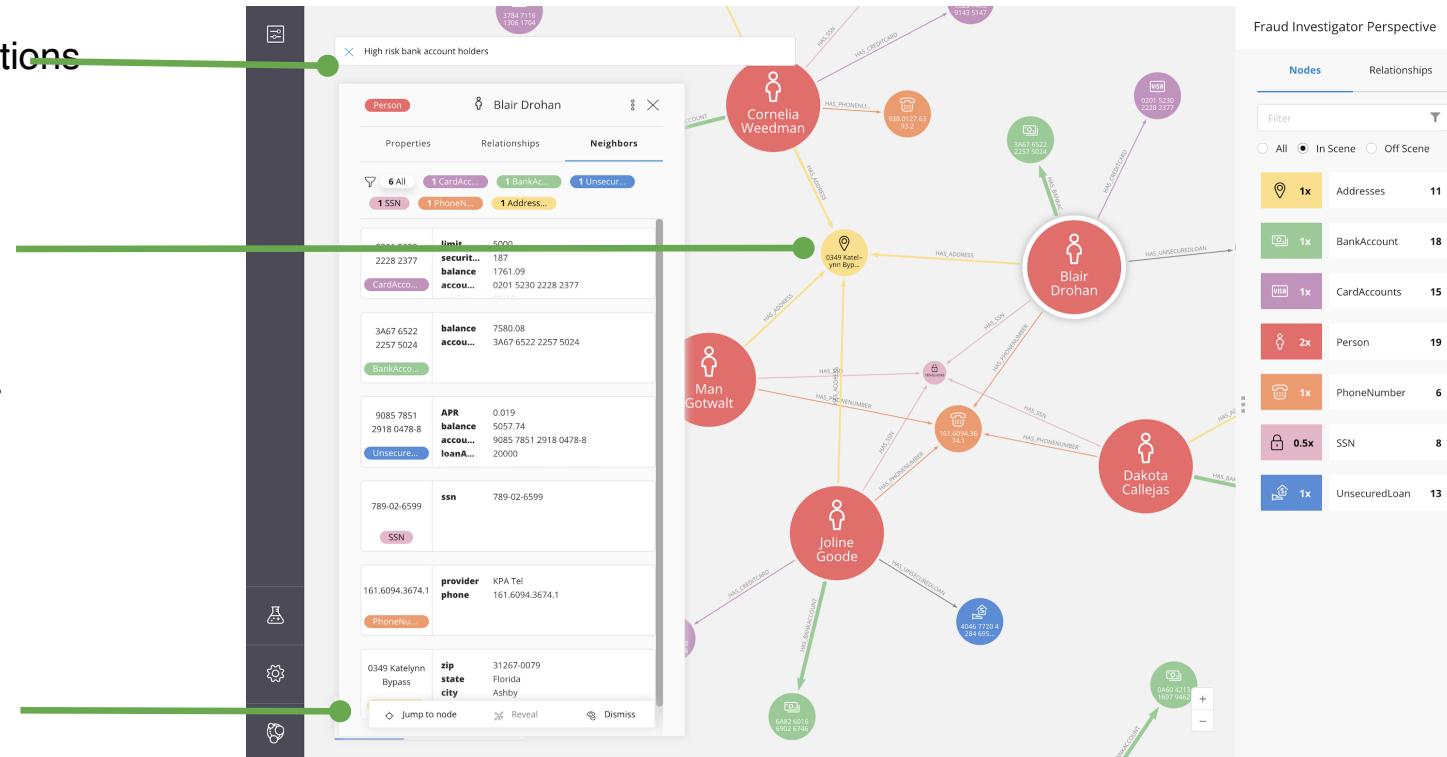
Search with type-ahead suggestions

Flexible Color, Size and Icon schemes

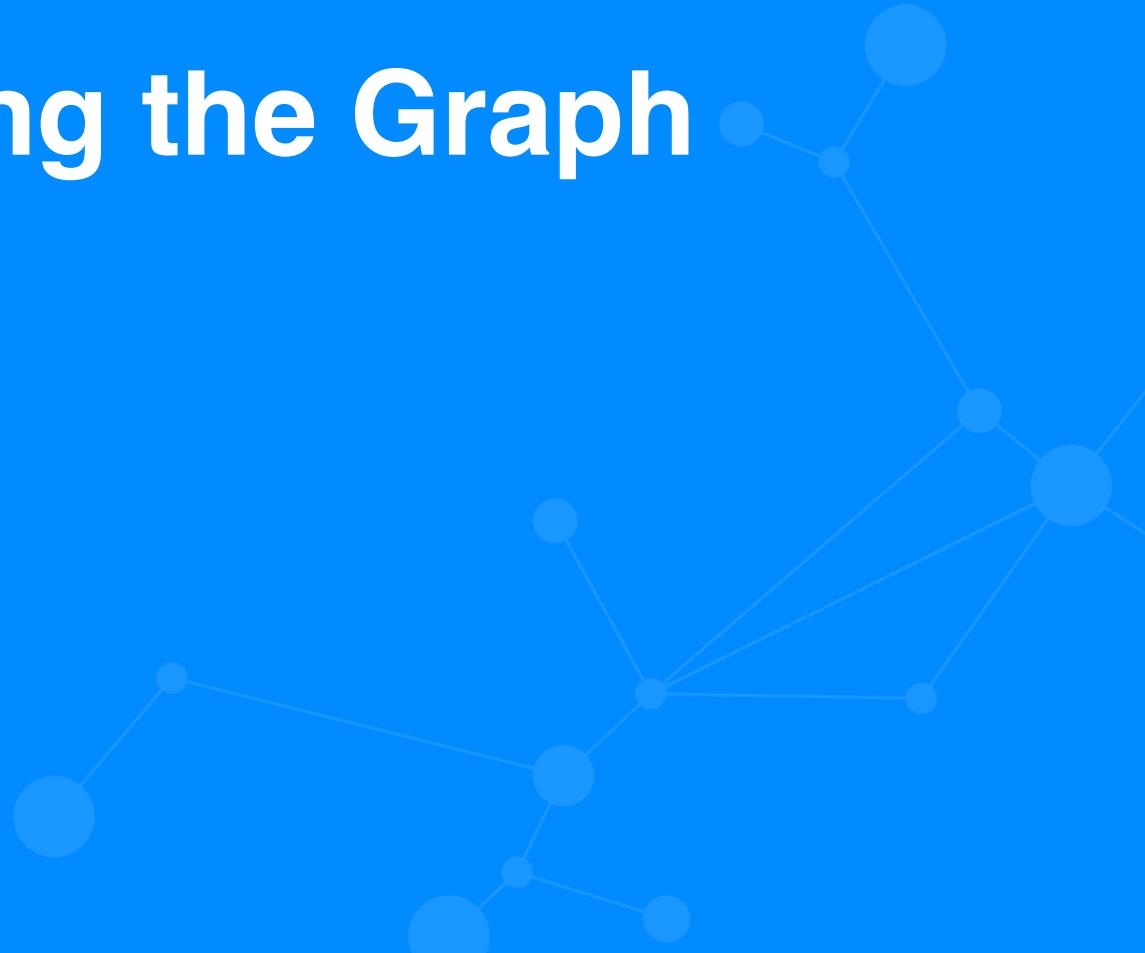
Visualize, Explore and Discover

Pan, Zoom and Select

Property Browser and editor

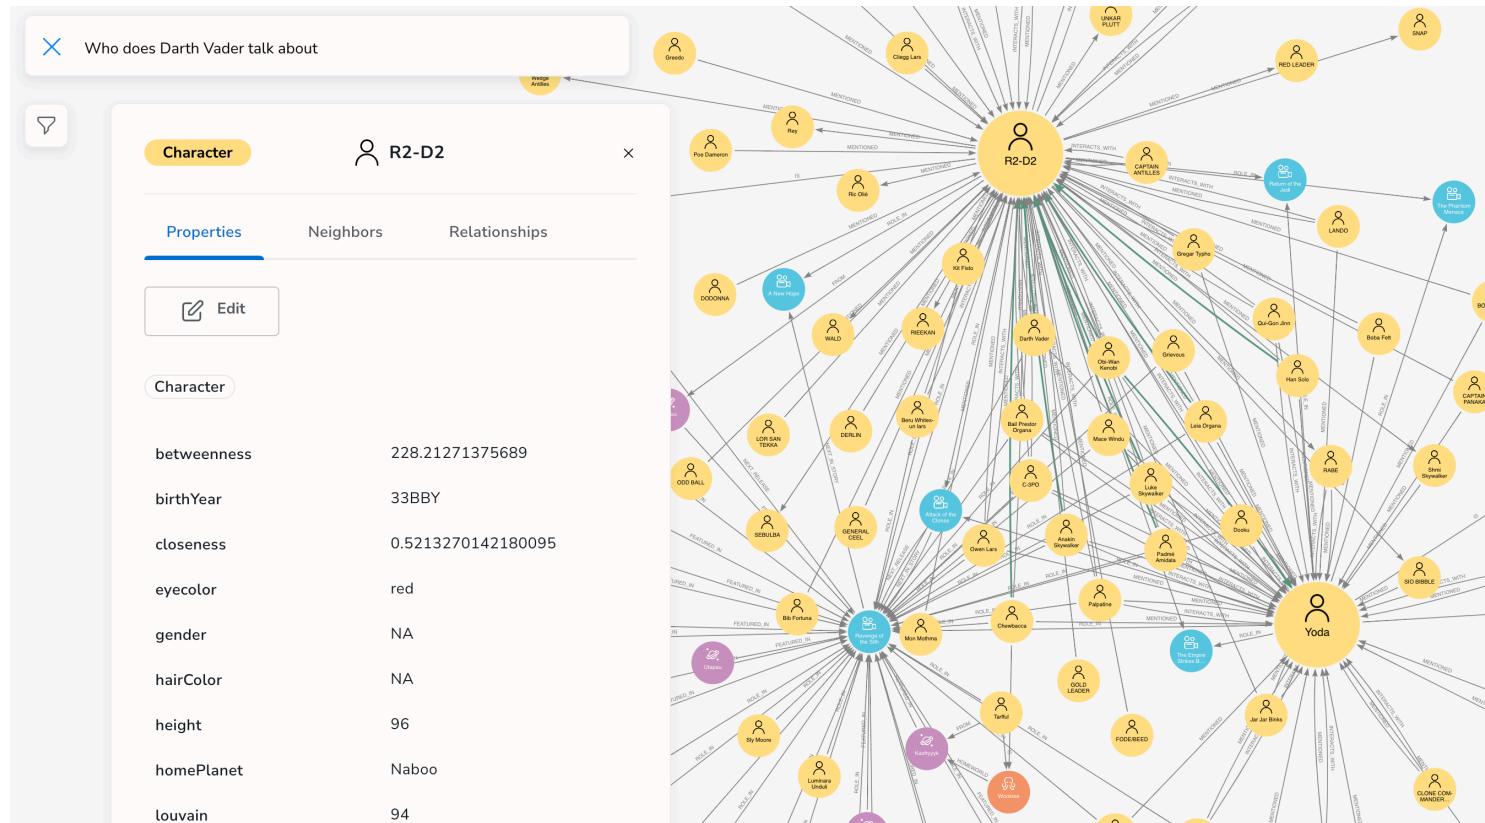


Lab 5 - Exploring the Graph with Bloom



Lab 5.1 - Exploring the Graph Using Bloom

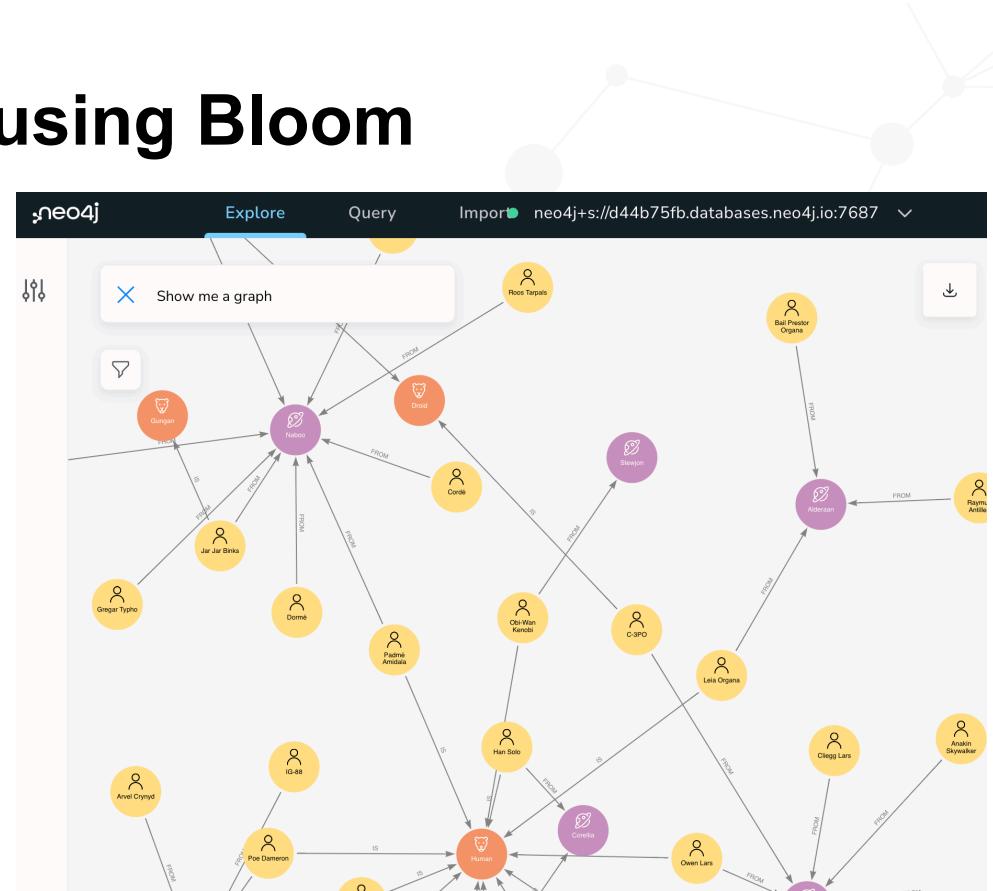
Explore the graph without writing code (unless you want to)



Lab 5.2 - Exploring Your DB using Bloom

Switch to Explore Tab

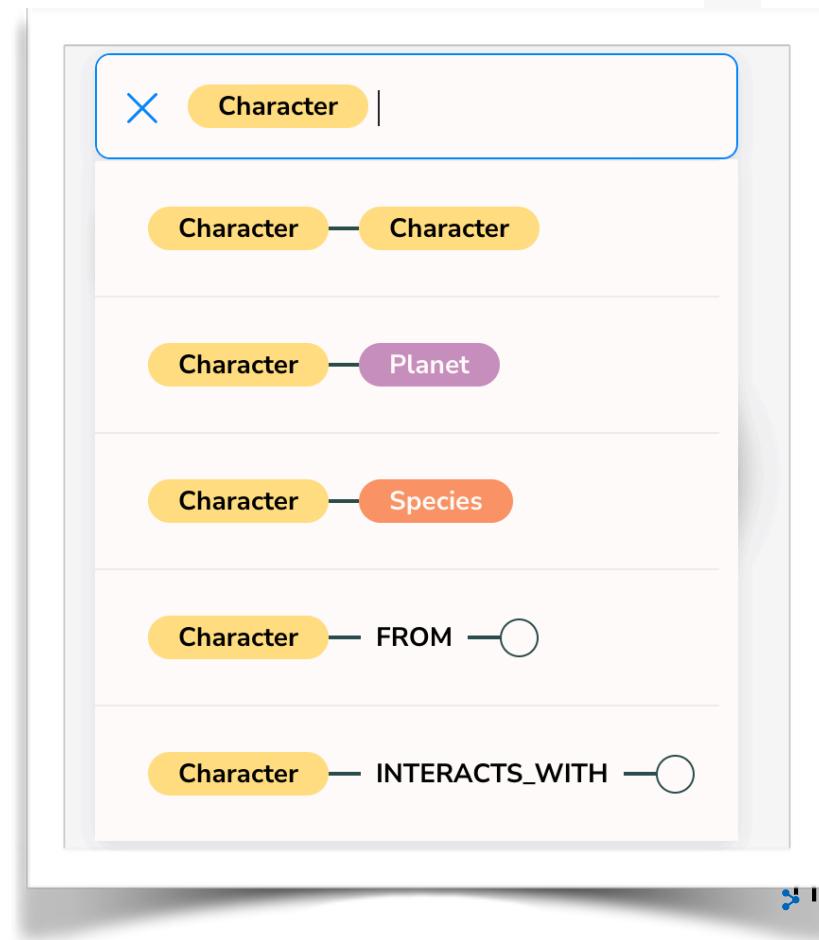
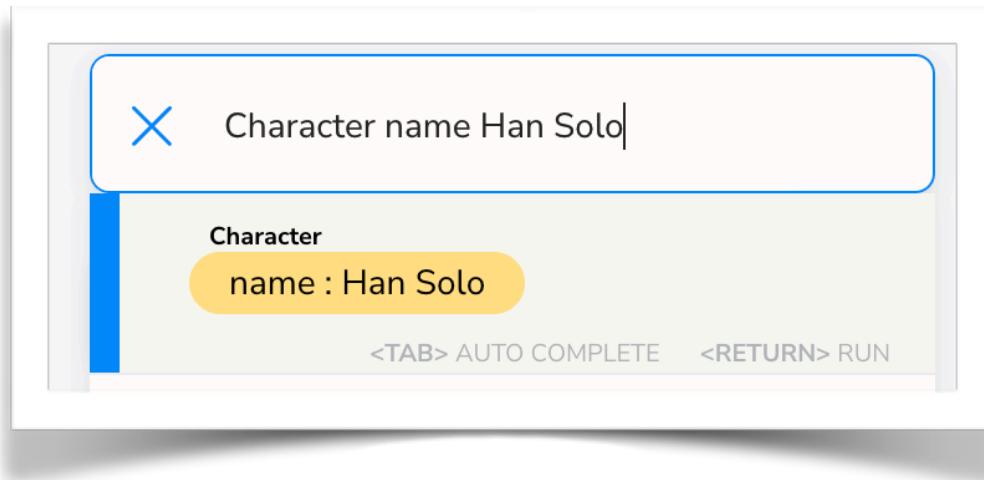
- Clicking on the Explore tab will invoke Neo4j Bloom.
- Bloom starts with a sample query to give an idea of the overall graph.



Lab 5.3 - Exploring Your DB using Bloom

Pattern Searching using NLP

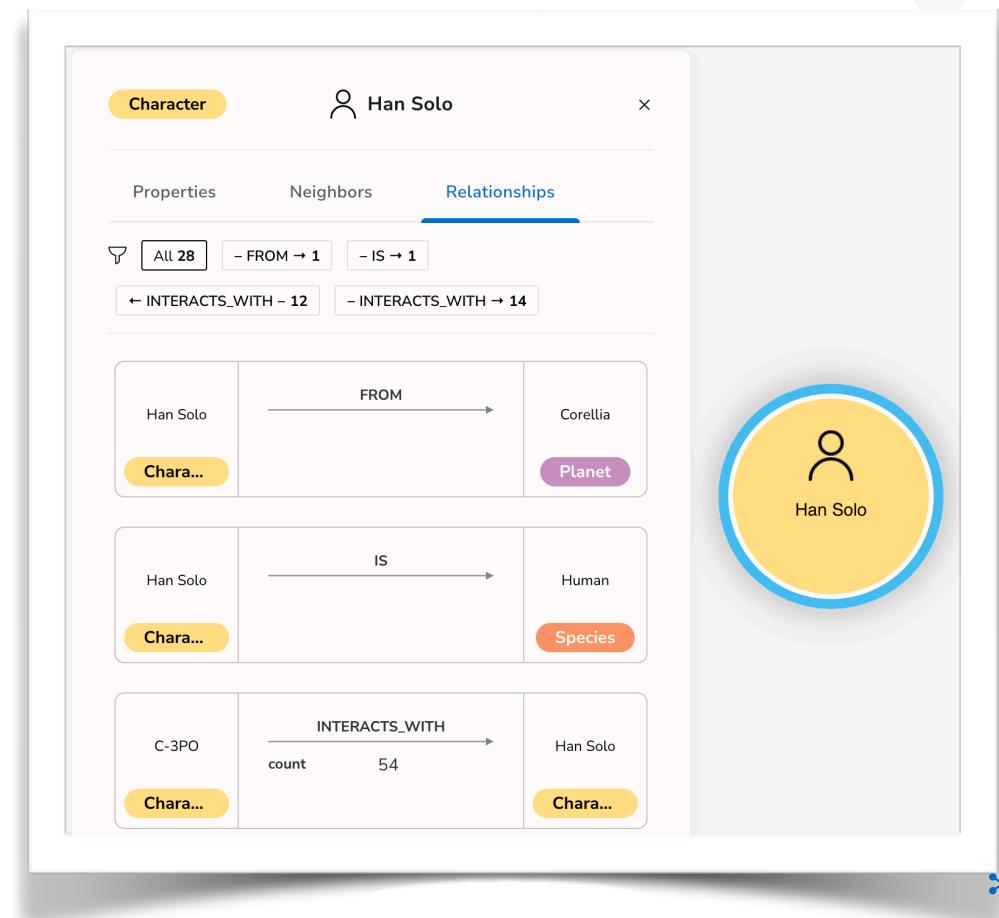
- Either type in an NLP phrase or build the pattern step by step



Lab 5.4 - Exploring Your DB using Bloom

Inspect Nodes to see properties and relationships

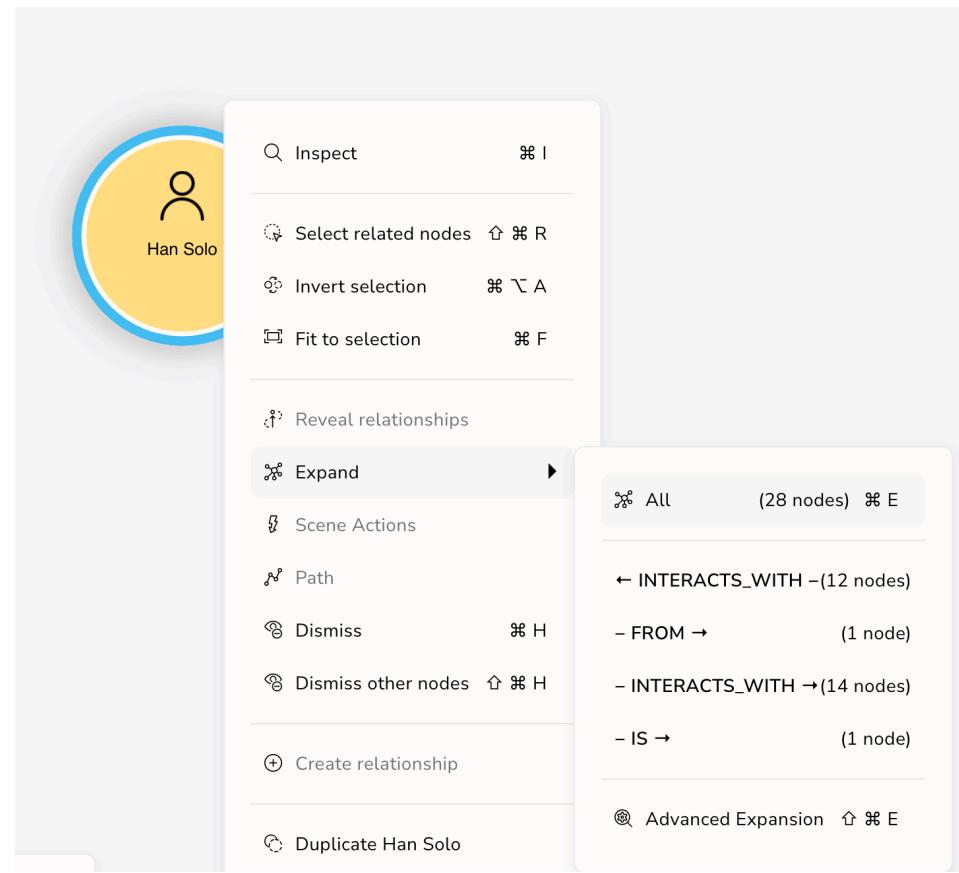
- Double-click on a node to inspect a node.
- Properties may be viewed and updated
- Related nodes can be viewed in the Neighbors tab or Relationships tab.



Lab 5.5 - Exploring Your DB using Bloom

Expanding Nodes to show related nodes

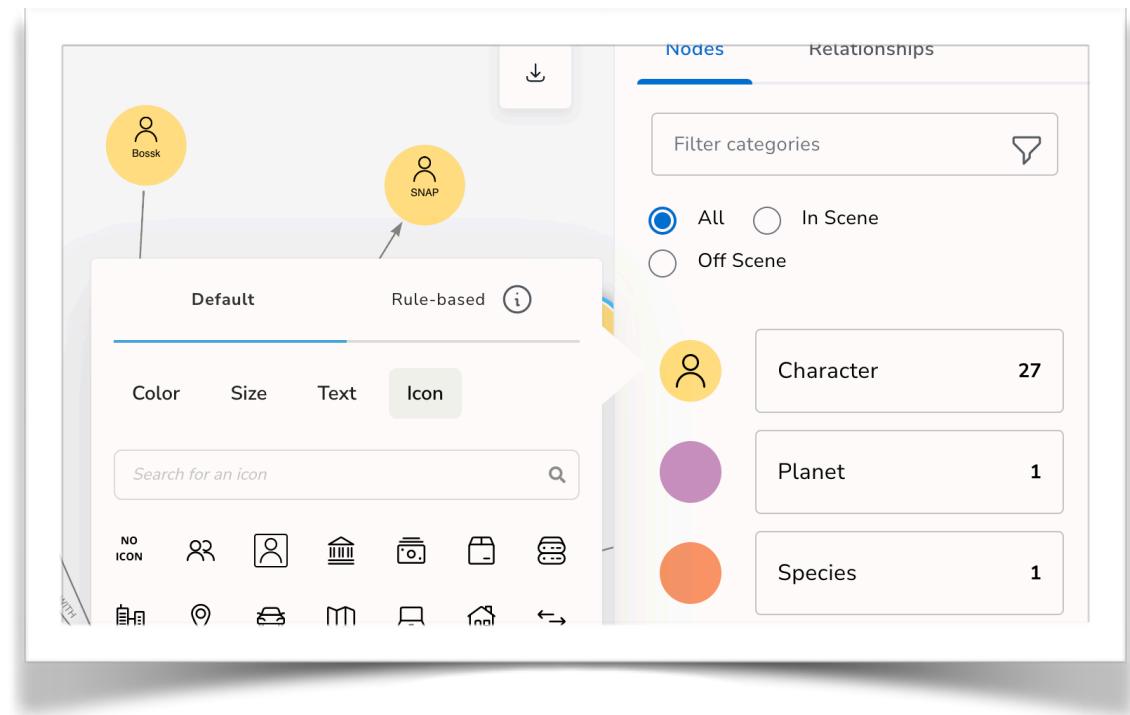
- Right click on a node to show pop-up menu for a node.
- Choose ‘Expand’ to show related nodes. Choose some or all relationships



Lab 5.6 - Exploring Your DB using Bloom

Customizing the Display

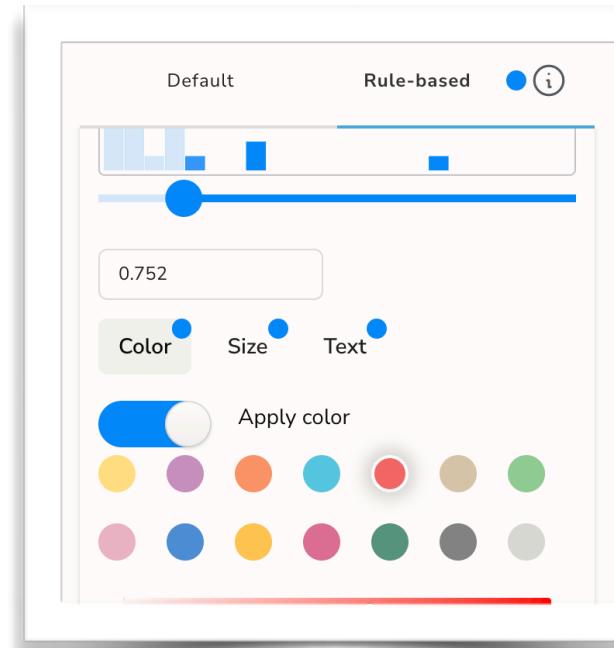
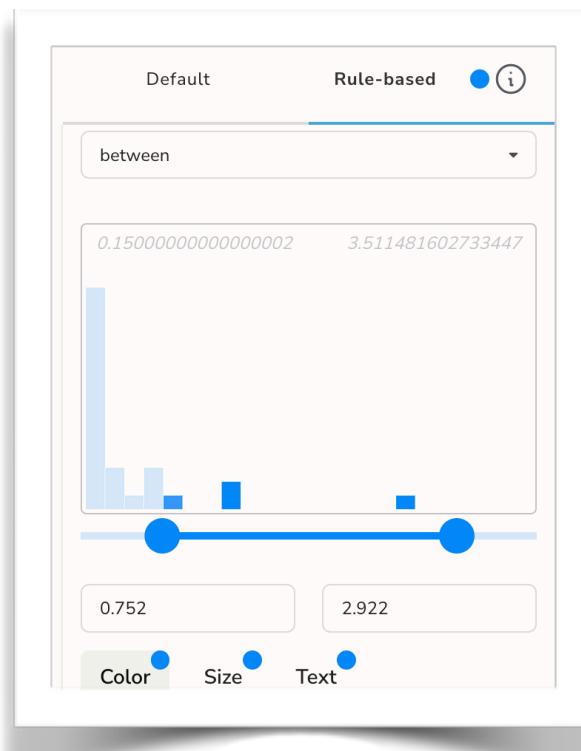
- Click on a Label to show display options.
- Color, size, text, and icons can all be selected.
- Rules can be used to highlight certain nodes



Lab 5.7- Exploring Your DB using Bloom

Changing the size and color of Characters with high PageRank scores

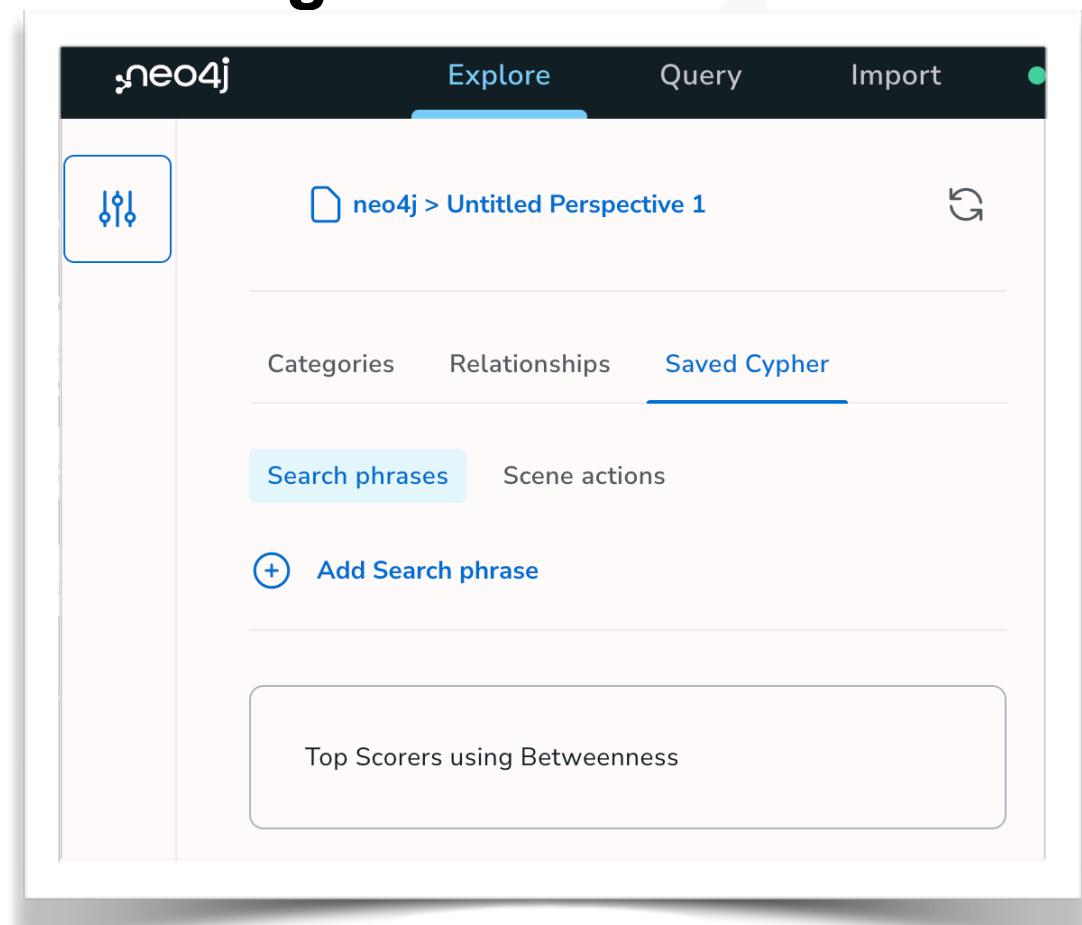
- Click on the Icon next to the Character Label
- Select Rule-based tab
- Choose the page rank property
- Choose single values and greater than or between.
- Use the slider to specify the values to use
- Choose a color and/or size for matching nodes



Lab 5 - Exploring Your DB using Bloom

Using Saved Cypher

- Click on the Perspectives Icon (upper left corner)
- Click on Saved Cypher / Search Phrases
- Add a search phrase



Lab 5 - Exploring Your DB using Bloom

Using Saved Cypher

- Give your search phrase a name and description
- Cut/paste the Cypher query that will be used
- Execute the query by typing its name into the search box

Top Scorers using Betweenness Delete

Search phrase *

Description

Cypher query *

```
MATCH (c:Character)
with
c ORDER BY c.pageRank
LIMIT 10
match p= (c)--(:Character)
return p
```

Lab 5.8 - Parameterized Queries

Add a parameter to a Cypher statement

- Give your search phrase a name and description
- Add a parameter to the search phrase using \$
- Create a parameterized Cypher statement
 - cut/paste this:

```
MATCH path=(p:Planet {name: $planet})<-[FROM]-(c:Character)
RETURN path
```
- Define how Bloom should suggest parameter values

© 2022 Neo4j, Inc. All rights reserved.

The screenshot shows the Neo4j Bloom interface for defining search configurations. The configuration is named "Characters from \$planet".
- **Search phrase**: "Characters from \$planet".
- **Description**: "Show Characters from a given planet".
- **Cypher query**:

```
MATCH path=(p:Planet {name:$planet})<-[FROM]-
(c:Character)
RETURN path
```

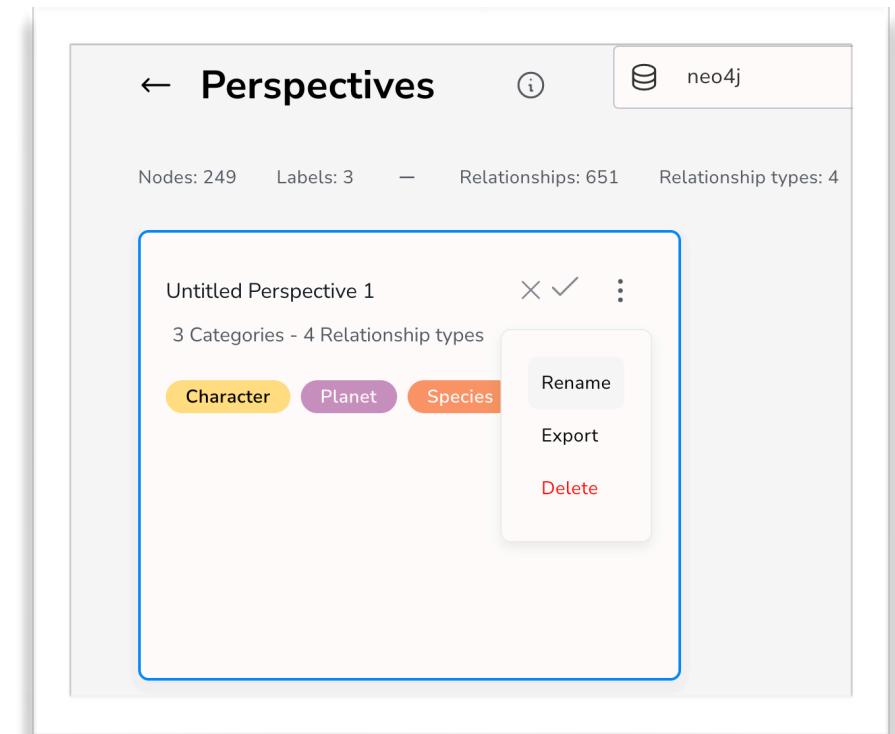

- **Parameter \$planet**:

- Data type:** String
- Search suggestions:** Label-key (with "Planet" selected)
- Value suggestions:** name (with "name" selected)

Lab 5.9- Export Perspectives

Use JSON to persist & share perspectives

- Perspectives may be exported to JSON files for persistence and versioning
- Export your perspective to a JSON file (optional)



Lab 5.10- Import Perspectives

Import JSON to use a different perspective

- Perspectives may be imported from saved JSON files.
- Multiple JSON files may be uploaded for easily jumping from one perspective to another.
- Enterprise versions of Bloom have additional methods for sharing perspectives amongst team members
- Example JSON file is https://github.com/mquinz/star_wars_demo/raw/master/bloom/demo_perspective.json

© 2022 Neo4j, Inc. All rights reserved.

The screenshot shows the Neo4j Bloom interface with the following details:

- Perspectives:** The main title bar shows "Perspectives".
- Database:** neo4j
- Nodes:** 249
- Labels:** 3
- Relationships:** 651
- Relationship types:** 4
- Neo4j version:** 5.4-aura
- Import:** A blue button in the top right corner.
- Untitled Perspective 12345:** Contains 3 Categories - 4 Relationship types. Categories shown: Character (yellow), Planet (purple), Species (orange).
- Simple Example for Training:** Contains 5 Categories - 4 Relationship types. Categories shown: Character (yellow), Planet (purple), Species (orange), Starship (blue), Vehicle (green).

Lab 6 - Using Python Notebooks



Lab 6.1 - Neo4j Using Python Notebook

Use official Neo4j Drivers

- Neo4j has 3 Python Drivers
- pip install neo4j

Execute Cypher read query and export to DataFrame

In [34]:

```
# load any parameters
params = { 'name': 'Han Solo' }

query = """MATCH (p:Character {name: $name})-[i:INTERACTS_WITH]-(p2:Character)
    RETURN p2.name AS otherName, i.count as interactionCount
    ORDER BY interactionCount DESC LIMIT 10
"""

# run query to get results
result = conn.query(query, parameters=params)

# create a dataframe from Cypher Resultset
dtf_data = pd.DataFrame([dict(_) for _ in result])
dtf_data.head(3)
```

Out [34]:

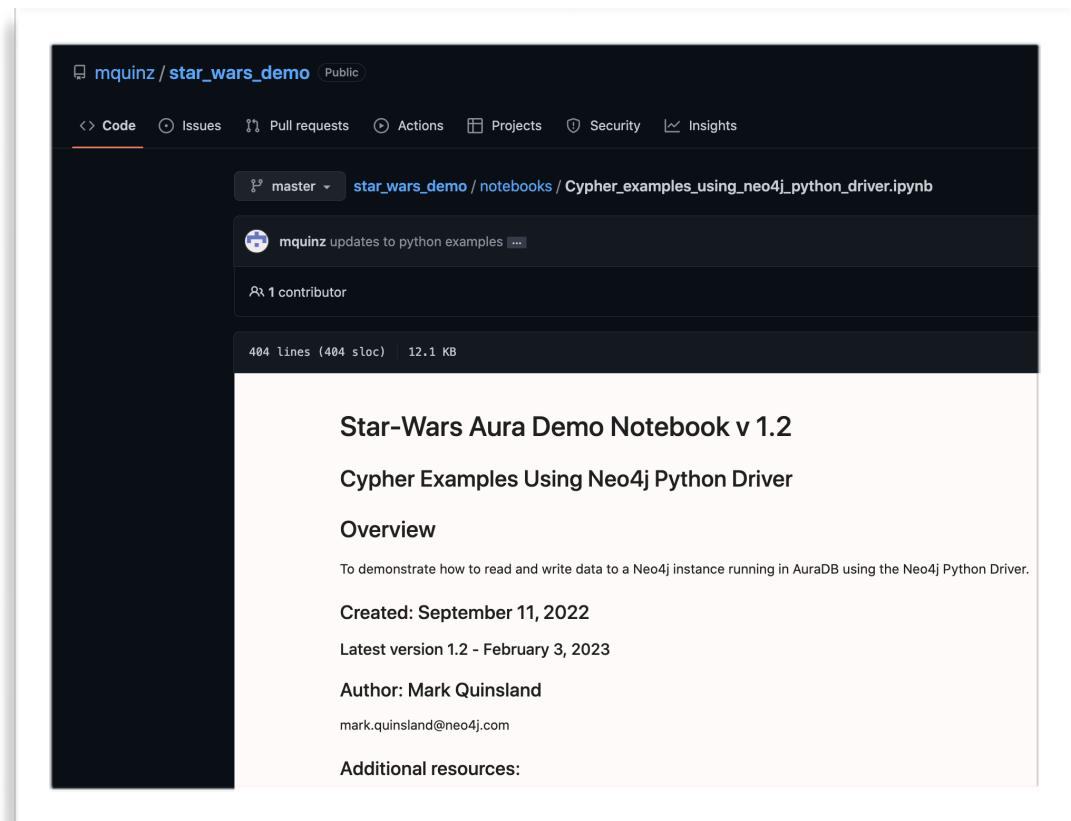
	otherName	interactionCount
0	Chewbacca	77
1	Leia Organa	69
2	C-3PO	54

https://github.com/mquinz/star_wars_demo/blob/master/notebooks/Cypher_examples_using_neo4j_python_driver.ipynb

Lab 6.2 - Neo4j Using Python Notebook

Download Example

- https://github.com/mquinz/star_wars_demo/blob/master/notebooks/Cypher_examples_using_neo4j_python_driver.ipynb



Lab 6.3 - Neo4j Using Python Notebook

Execute Query and Process Results

```
# load any parameters using a dictionary

params = {'name': 'Han Solo' }

# cypher - who does Han Solo interact with the most?
query = """MATCH (p:Character {name: $name})-[i:INTERACTS_WITH]-(p2:Character)
            RETURN p2.name AS otherName, i.count as interactionCount
            ORDER BY interactionCount DESC LIMIT 10

"""

result = conn.query(query, parameters=params)

for row in result:
    print (row['otherName'], row['interactionCount'])
```



Lab 6.4 - Neo4j Using Python Notebook

Load Query Results to DataFrame

```
# load any parameters

params = {'name': 'Han Solo' }

query = """MATCH (p:Character {name: $name})-[i:INTERACTS_WITH]-(p2:Character)
    RETURN p2.name AS otherName, i.count as interactionCount
    ORDER BY interactionCount DESC LIMIT 10
"""

# run query to get results
result = conn.query(query, parameters=params)

# create a dataframe from Cypher Resultset
dtf_data = pd.DataFrame([dict(_) for _ in result])
dtf_data.head(3)
```

	otherName	interactionCount
0	Chewbacca	77
1	Leia Organa	69
2	C-3PO	54

Lab 6.5 - Neo4j Using Python Notebook

Insert Nodes and Relationships Using Batches

- Perspectives may be exported to JSON files for persistence and versioning
- Export your perspective to a JSON file (optional)

```
params = { }

# create a batch of dictionaries - each item will be processed individually within the batch
bunch_of_rows = [
    {'name': 'Mark', 'gender': 'male', 'planet': 'Endor', 'jedi': 'true'},
    {'name': 'Abby', 'gender': 'female', 'planet': 'Endor', 'jedi': 'false'},
    {'name': 'Grokster', 'gender': 'male', 'planet': 'Alderaan', 'jedi': 'false'}
]

params ['rows'] = bunch_of_rows

query = """

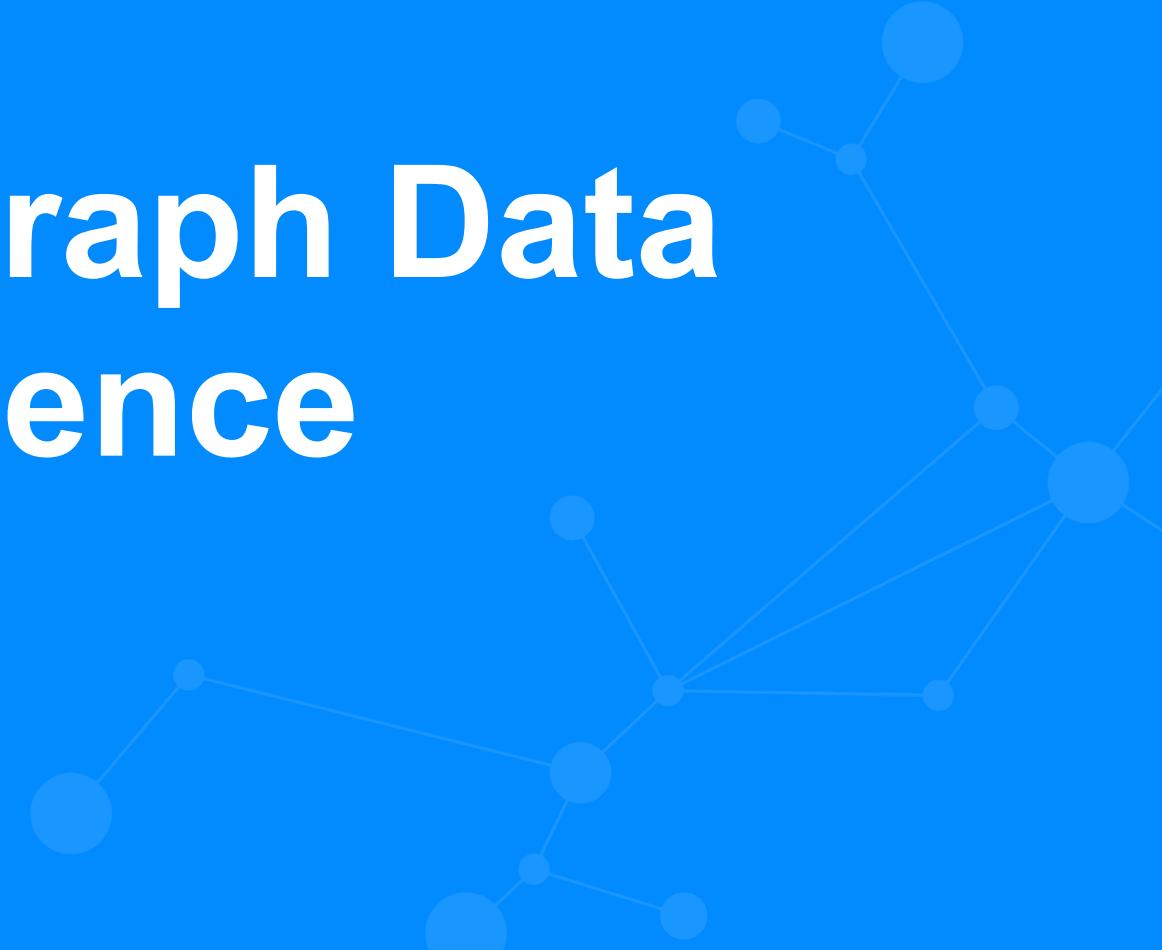
// use UNWIND to process each item in the array
UNWIND $rows AS row
MERGE (u:Character {name: row.name})
SET u.gender = row.gender,
    u.jedi = toBoolean(row.jedi)

// now create a relationship from the Character to their home planet

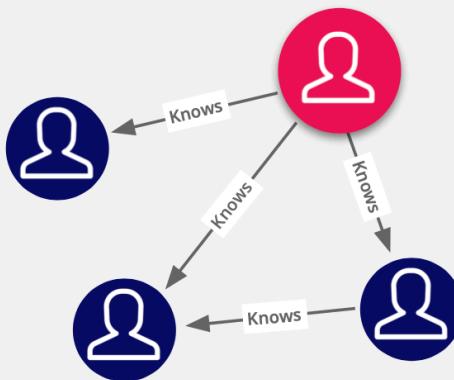
WITH u, row.planet as homePlanet
    // find the planet and create a relationship to it
    MATCH (p:Planet)
        WHERE p.name = homePlanet
    MERGE (u)-[:FROM]->(p)
return count(*) as total
"""

# run query to get results
result = conn.query(query, parameters=params)
print (result)
```

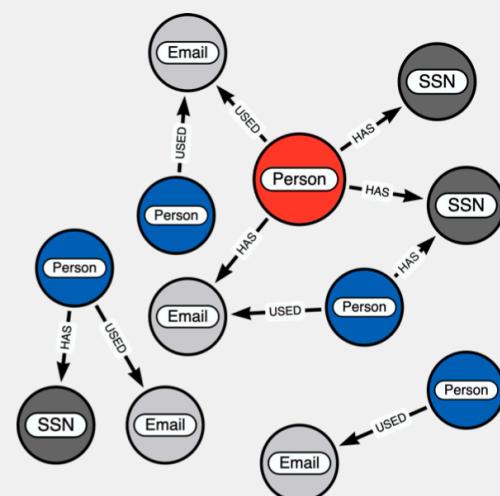
Neo4j Graph Data Science



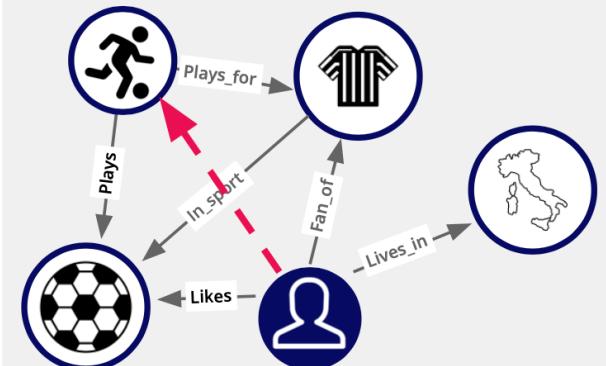
Use Graph Data Science to Answer:



What's important?



What's unusual?



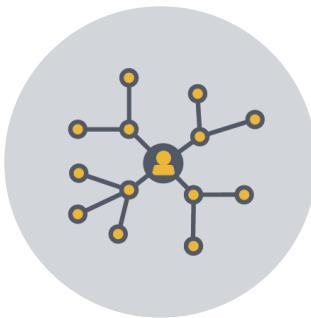
What's next?

Graph Algorithm Categories



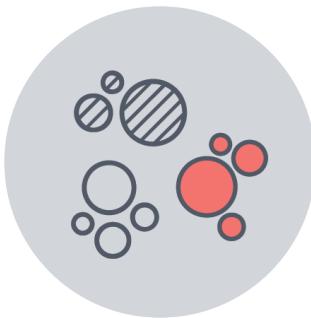
Pathfinding and Search

Finds optimal paths or evaluates route availability and quality.



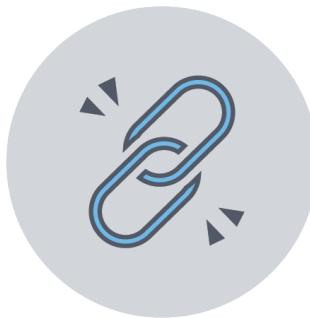
Centrality

Determines the importance of distinct nodes in the network.



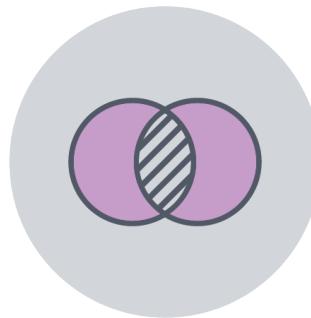
Community Detection

Detects group clustering or partition.



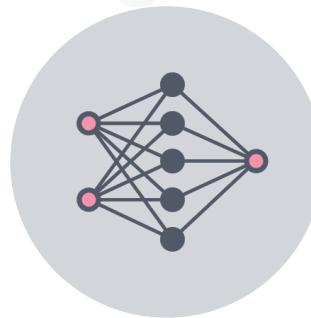
Heuristic Link Prediction

Estimates the likelihood of nodes forming a future relationship.



Similarity

Evaluates how alike nodes are by neighbors and relationships.



Embeddings

Learns graph topology to reduce dimensionality for ML

Graph Centrality Algorithms Assist in Analysis

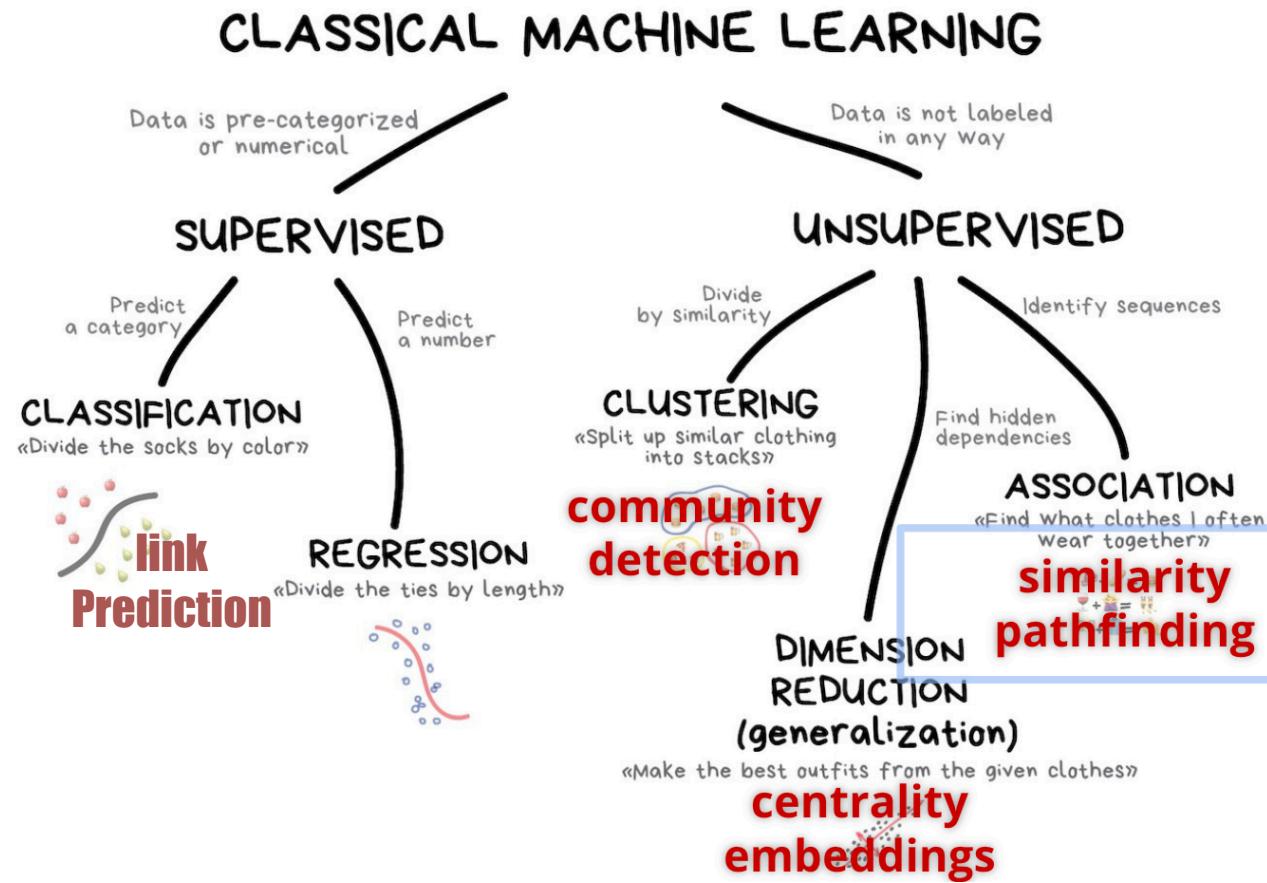


- Closeness algorithms identify:
 - Highly-connected nodes
 - Built-in redundancy
- Betweenness algorithms identify:
 - Strategically-connected nodes
 - Single points of failure
 - Potentially vulnerable nodes

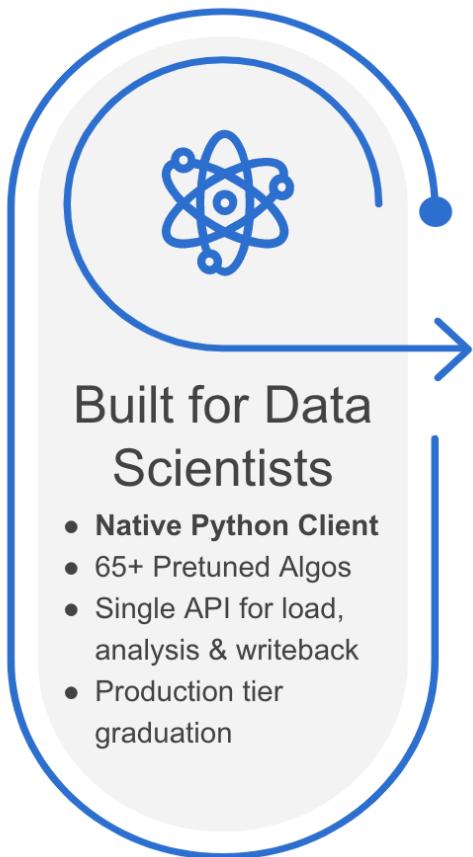
Neo4j GDS - Part of your ML Pipeline

Machine learning can be

- **unsupervised:** unlabelled, identifying patterns, or
- **supervised:** labeled data, predicting values



Neo4j Graph Data Science - End User Experience



Connect to Graph Data Science (GDS)

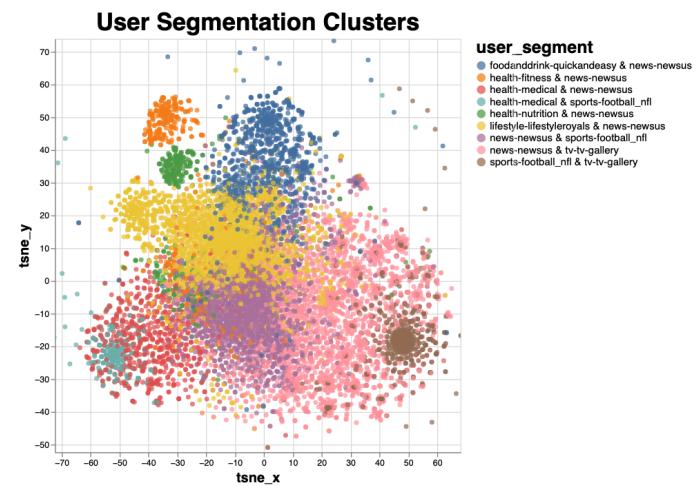
```
In [3]: from graphdatascience import GraphDataScience  
  
# Use Neo4j URI and credentials according to your setup  
gds = GraphDataScience(HOST, auth=(USERNAME, PASSWORD), aura_ds=True)
```

Apply GDS for Automated Graph Feature Engineering

FastRP transforms the graph of interconnected user activity into embedding features for EDA & ML applications

```
In [4]: g, _ = gds.graph.project('proj', ['User', 'News'], {'CLICKED': {'orientation': 'UNDIRECTED'}})  
gds.fastRP.mutate(g, mutateProperty='segmentEmbedding', embeddingDimension=128, randomSeed=7474)  
gds.graph.writeNodeProperties(g, ['segmentEmbedding'], ['User'])
```

```
Out[4]: writeMillis 1212  
graphName proj  
nodeProperties [segmentEmbedding]  
propertiesWritten 750434  
Name: 0, dtype: object
```



Let's Do Something Amazing Together

neo4j